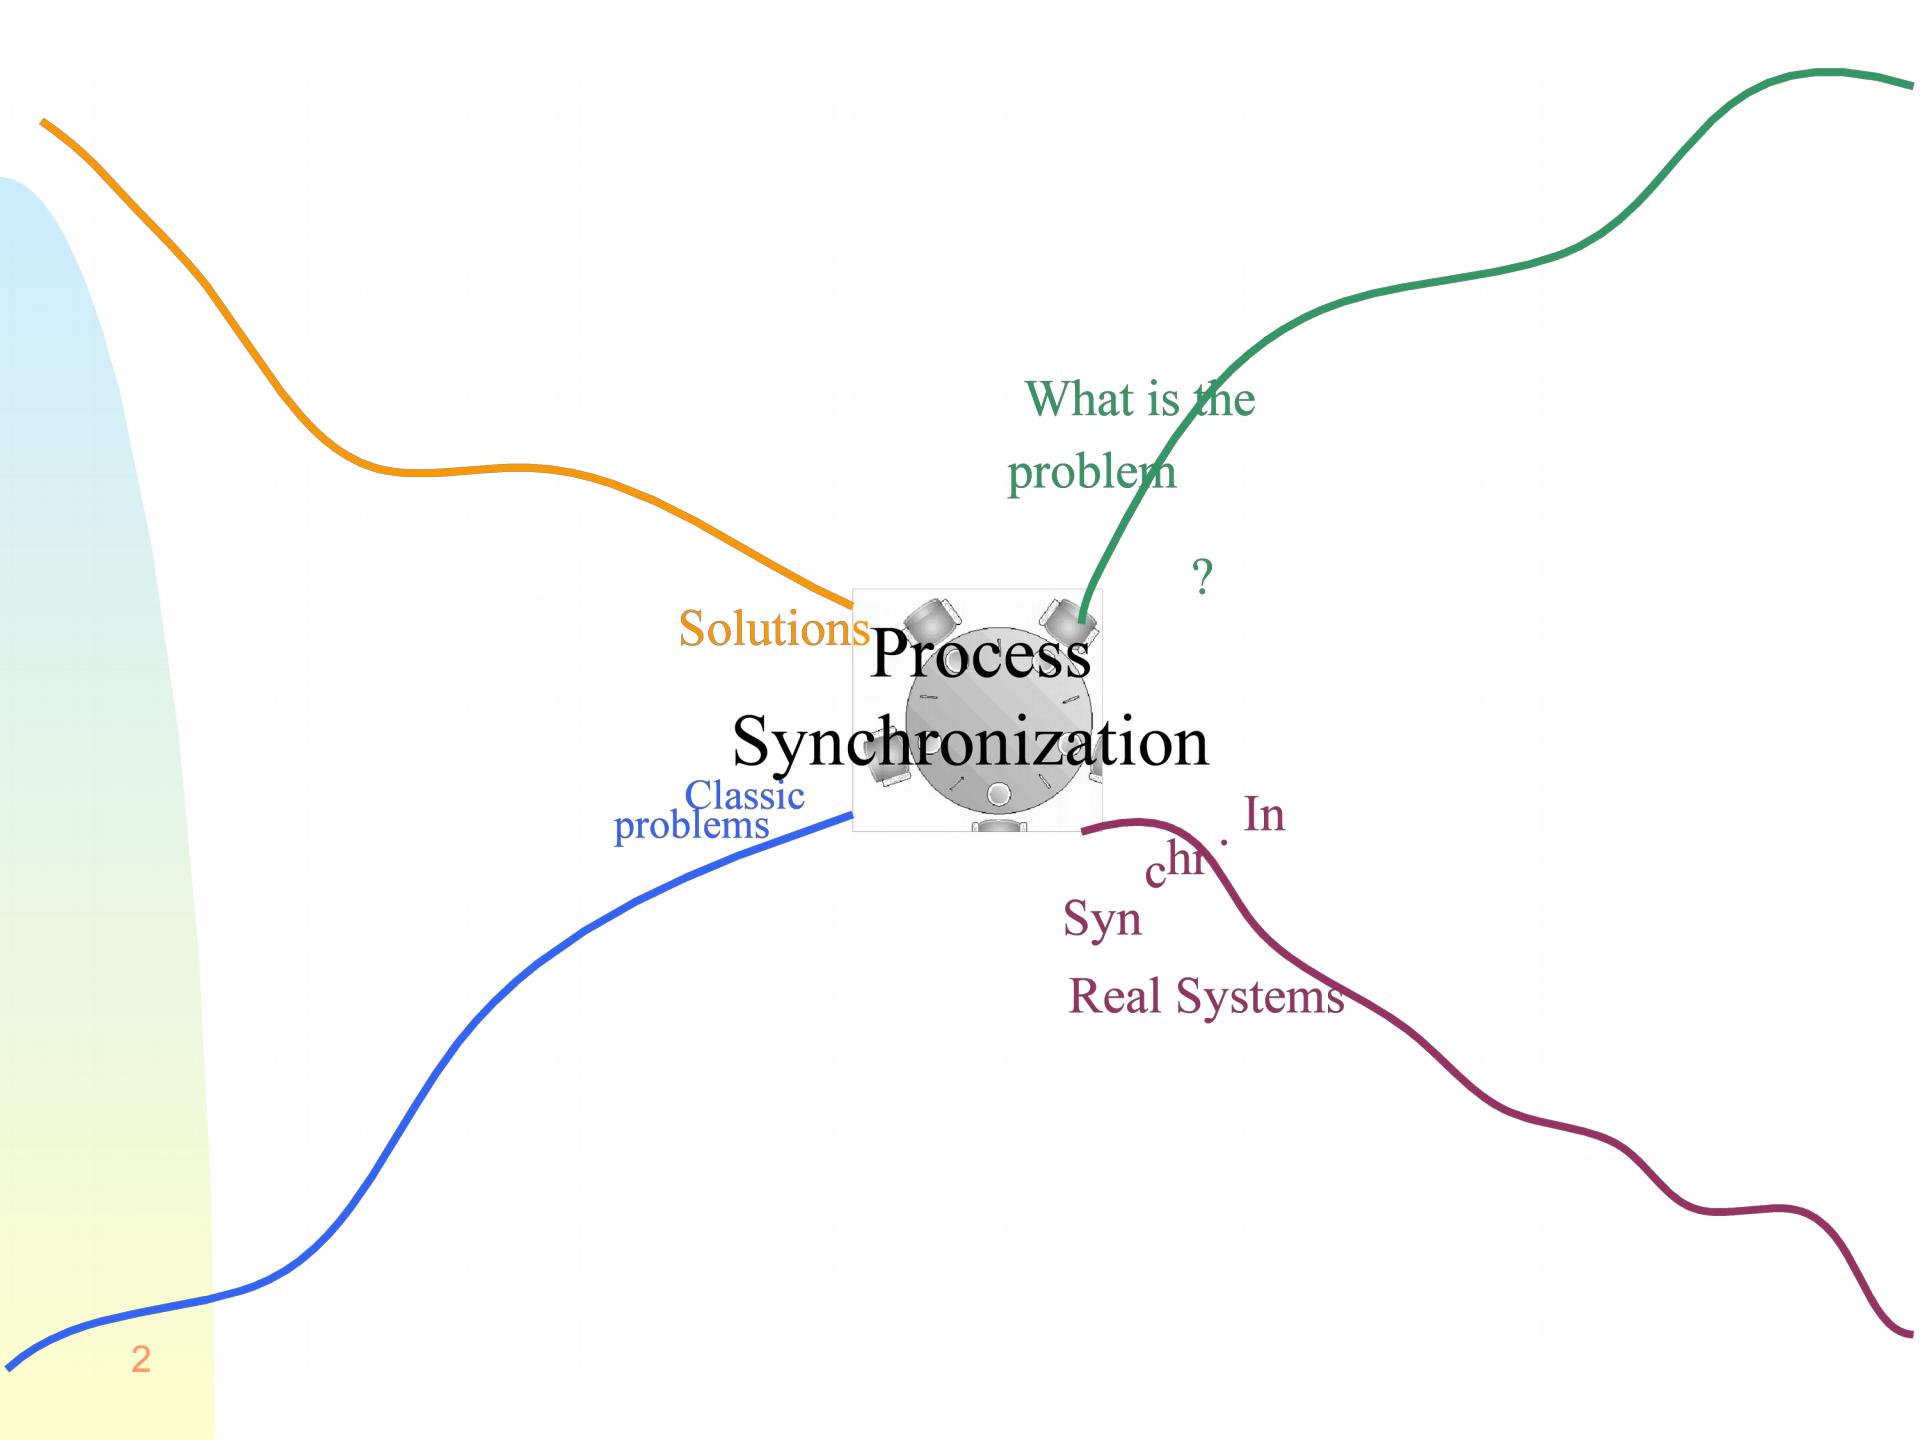
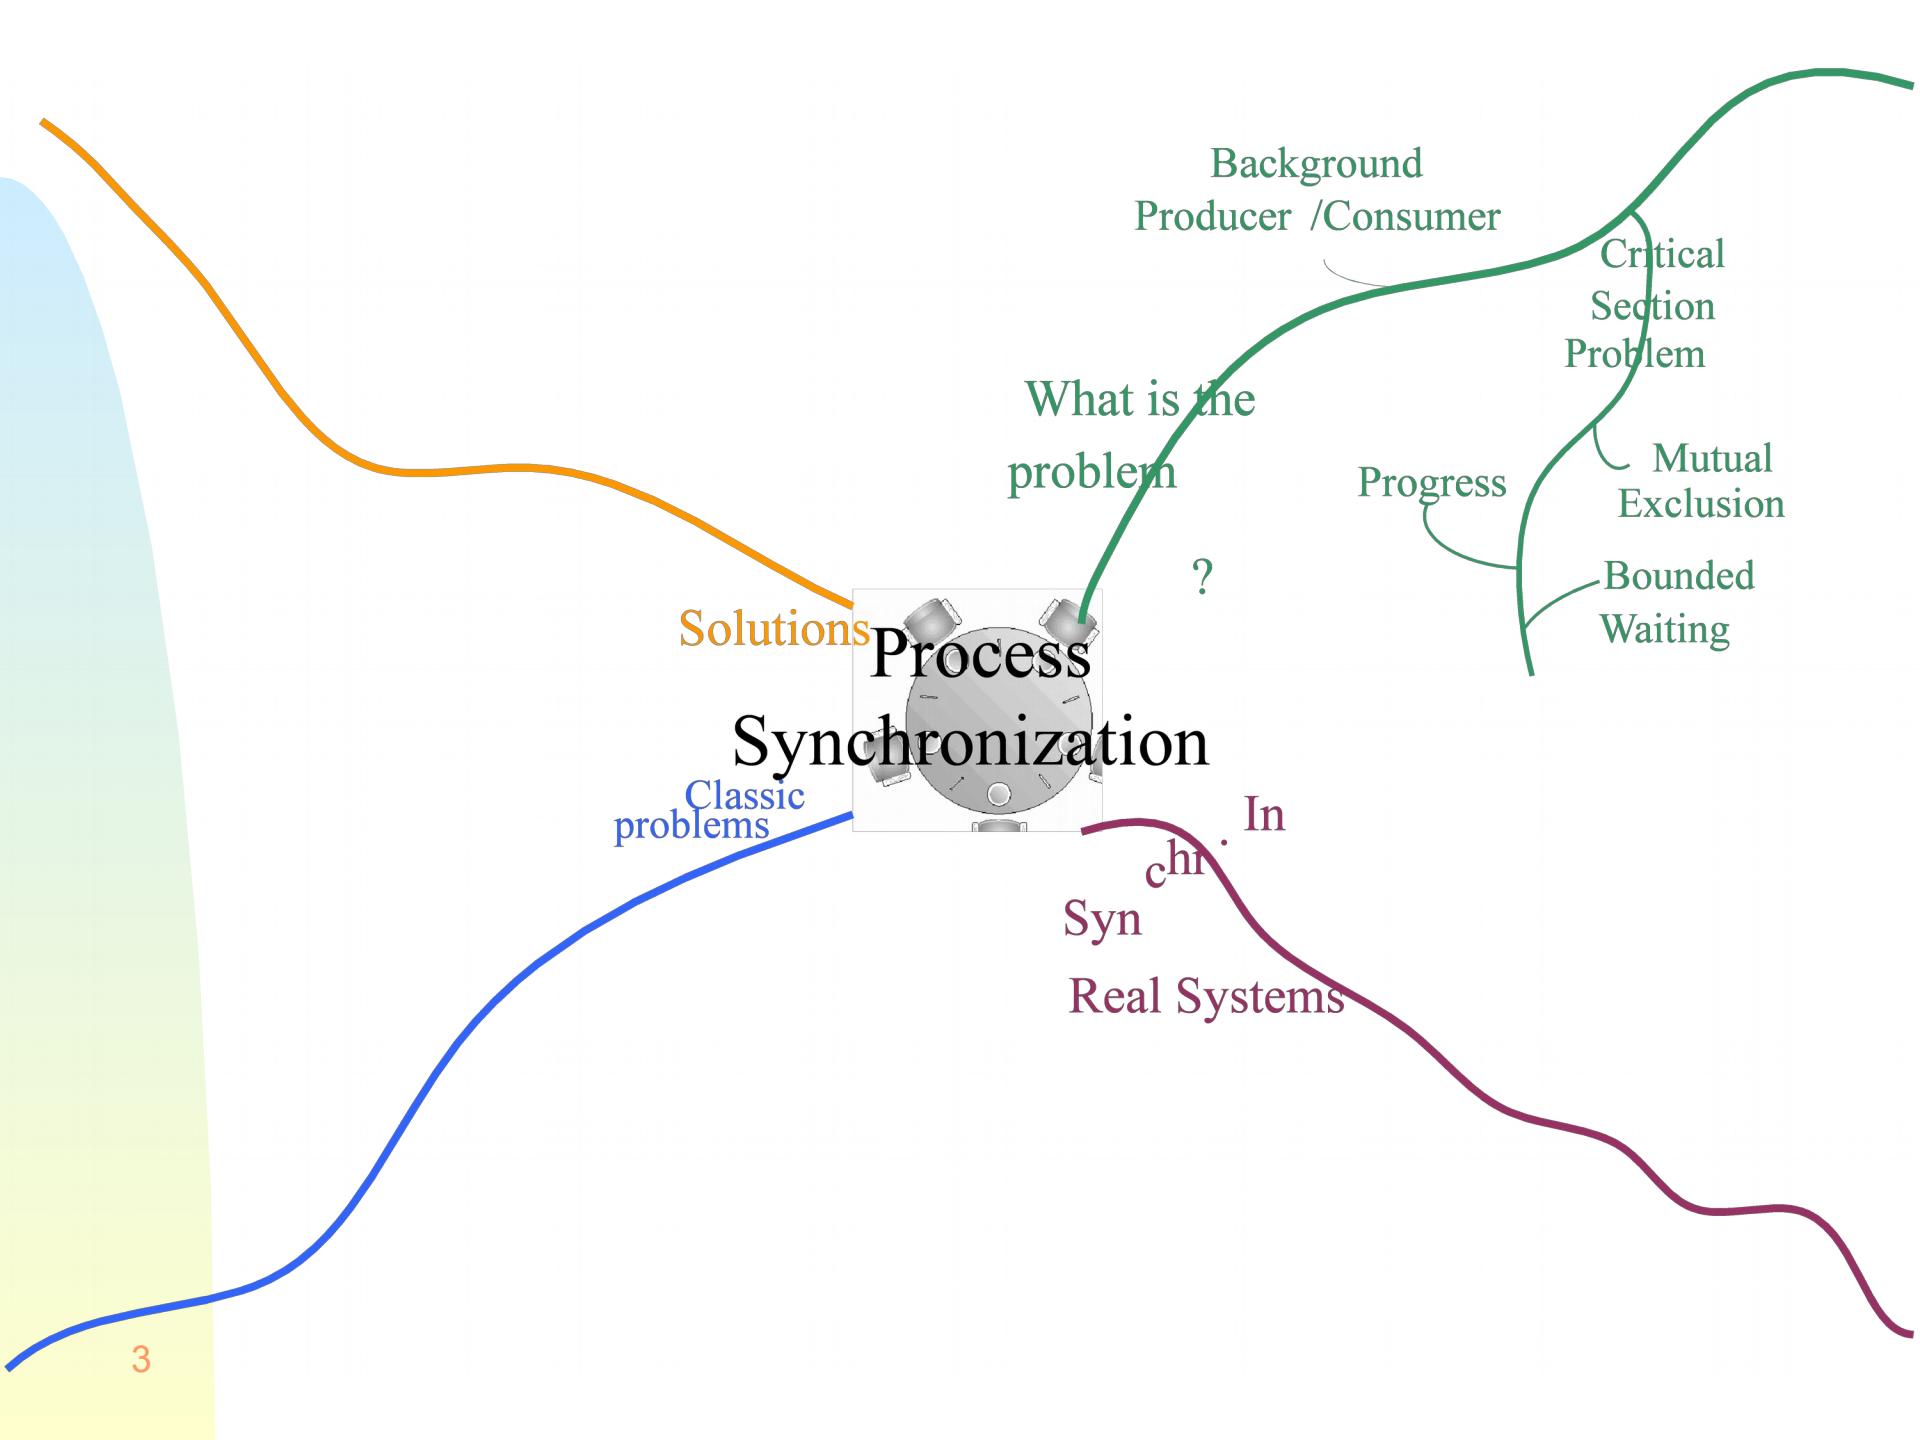


# **Module 5: Process Synchronization**

---

**Reading: Chapter 6**





# Problems with concurrence = parallelism

- Concurrent threads sometimes share data (files and common memory) and resources
  - These are cooperative tasks
- If access is not controlled, the result of the program execution may depend on the order of interleaving of instruction execution (non-deterministic).
- A program can give different results from one execution to another and sometimes undesirable results.

# Is there a problem?

Consider the following code, in which a child is sharing memory with the parent that created it.

What is the output produced by this code?

```
a = 2;  
b = 3;  
pid = clone();  
if (pid < 0)  
    // error handling  
else if (pid == 0)  
{  
    // child thread  
    b = 4;  
    c = a+b;  
    printf("%d\n", c);  
}  
else  
{  
    // parent thread  
    a = 0;  
    c = a+b;  
    printf("%d\n", c);  
}
```

# Shared Data Structure Example

OK, that was silly, what about a more real-life example?

*Very simplified queue:*

```
enqueue (data) :  
    qData[inPos] = data;  
    inPos++;  
  
dequeue () :  
    data = qData[outPos];  
    outPos++;
```

**What happens if several processes/threads try to use the queue functions concurrently?**

# Producer-Consumer Example

The Producer executes

```
while (true)
{
    while (count == BUFFER_SIZE)
        ; // do nothing
    // produce an item and put in
    // nextProduced
    buffer[inIx] = nextProduced;
    inIx = (inIx + 1) % BUFFER_SIZE;
    count++;
}
```

# Producer-Consumer Example

The Consumer executes

```
while (true)
{
    while (count == 0)
        ; // do nothing
    nextConsumed = buffer[outIx];
    outIx = (outIx + 1) % BUFFER_SIZE;
    count--;
    // consume the item in nextConsumed
}
```

So, does it work? Why?

- Lets see...

# Producer-Consumer Example

- count++ could be implemented as

```
register1 = count  
register1 = register1 + 1  
count = register1
```

- count-- could be implemented as

```
register2 = count  
register2 = register2 - 1  
count = register2
```

- Consider this execution interleaving:

S0: producer execute `register1 = count {register1 = 5}`  
S1: producer execute `register1 = register1+1 {register1 = 6}`  
S2: consumer execute `register2 = count {register2 = 5}`  
S3: consumer execute `register2 = register2-1 {register2 = 4}`  
S4: producer execute `count = register1 {count = 6 }`  
S5: consumer execute `count = register2 {count = 4}`

# Summary of the problem

- Concurrent processes (or threads) often need to share data (maintained either in shared memory or files) and resources
- If there is no controlled access to shared data, some processes will obtain an inconsistent view of this data
- The results of actions performed by concurrent processes will then depend on the order in which their execution is interleaved – race condition
- Let us abstract the danger of concurrent modification of shared variables into the critical-section problem.

# Critical-Section Problem

- The piece of code modifying the shared variables where a thread/process needs exclusive access to guarantee consistency is called critical section.
- The general structure of each thread/process can be seen as follows:

```
while (true)
{
    entry_section
    critical section (CS)
    exit_section
    remainder section (RS)
}
```

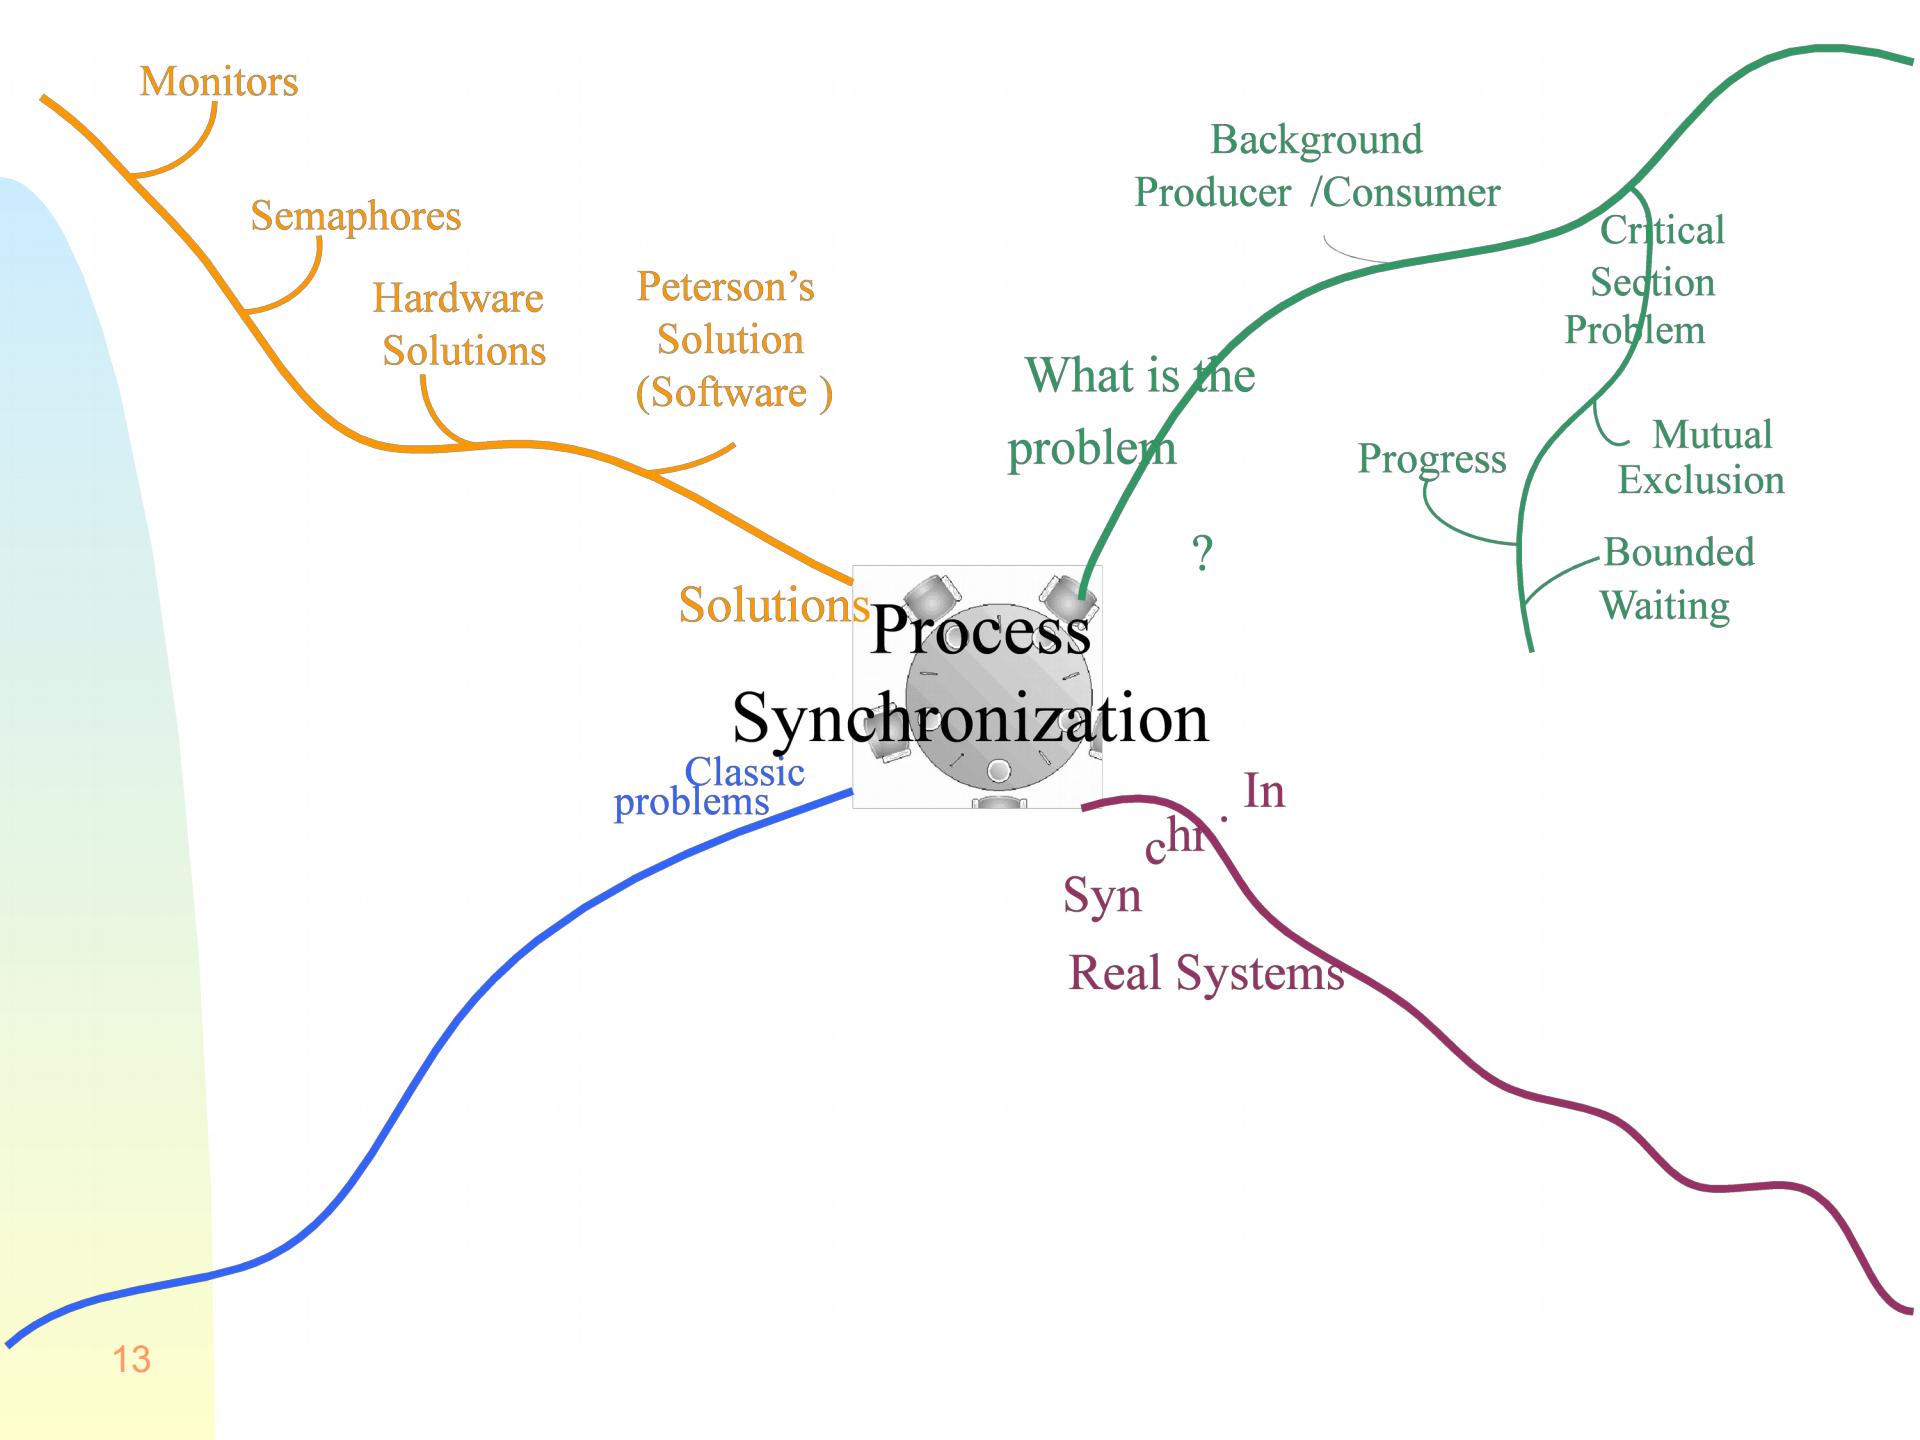
- Critical (CS) and remainder (RS) sections are given,
- We want to design entry and exit sections so that the following requirements are satisfied:

# Solution Requirements for CSP

1. **Mutual Exclusion** - If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If there exist some processes wishing to enter their CS and no process is in their CS, then one of them will eventually enter its critical section.
  - ♦ **No deadlock**
  - ♦ **Non interference** – if a process terminates in the RS, other processes should still be able to access the CS.
  - ♦ **Assume that a thread/process always exits its CS.**
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
  - ♦ **No famine.**

## Assumptions

- Assume that each process executes at a nonzero speed
- No assumption concerning relative speed of the  $N$  processes
- Many CPUs may be present but memory hardware prevents simultaneous access to the same memory location
- No assumption about the order of interleaved execution



# CSP Solutions



- Software solutions
  - **algorithms who's correctness does not rely on any other assumptions**
  - **Peterson's algorithm**
- Hardware solutions
  - **rely on some special machine instructions**
  - **testAndSet, xchng**
- OS provided solutions
  - **higher level primitives (implemented usually using the hw solutions) provided for convenience by the OS/library/language**
  - **semaphores, monitors**
- **All solutions are based on atomic access to memory: memory at a specific address can only be affected by one instruction at a time, an thus one thread/process at a time.**
- **More generally, all solutions are based on the existence of atomic instructions, that operate as basic CSs.**

# Software Solutions

Lets start with a simpler problem – only two tasks

- Two tasks,  $T_0$  and  $T_1$  (also  $T_i$  and  $T_j$ , where  $i \neq j$ )
- We want to design the entry and exit sections

```
while (true)
{
    entry_section
    critical section (CS)
    exit section
    remainder section (RS)
}
```

# Algorithm 1

Idea:

- Have a shared variable **turn** indicating whose turn it is now. It is initialized to 0 or 1, does not matter.
- Task  $T_i$  may enter the critical section if and only if  $\text{turn} = i$
- After exiting the critical section, **turn** is set to the other value, to let the other task gain access to its critical section
- $T_i$  can be in a **busy wait** if  $T_j$  is in its CS.

Task T0:

```
while(true)
{
    while(turn!=0) /*bw*/
    Critical Section
    turn=1;
    Remainder Section
}
```

Task T1:

```
while(true)
{
    while(turn!=1) /*bw*/
    Critical Section
    turn=0;
    Remainder Section
}
```

# Algorithm 1

Discussion:

- Does it ensure mutual exclusion?
  - Yes, at any moment, turn has only one value, and if a task  $T_i$  is in its critical section, then  $turn = i$
- Does it satisfy the bounded waiting requirement?
  - Yes, the tasks alternate
- Does it satisfy the progress requirement?
  - No, because it requires strict alternation of CSs
  - If a task requires its CS more often than the other, it cannot get it.

```
Task Ti:  
while(true)  
{  
    while(turn!=i) {};  
    Critical Section  
    turn=j;  
    Remainder Section  
}
```

# Algorithm 2

OK, Algorithm 1 was not that great, we do not need to alternate.

Idea:

- Lets introduce shared variables `flag[0]` and `flag[1]` to indicate whether task 0 and task 1 resp. wants to enter the CS
- Task i sets to true the `flag[i]` before trying to enter the CS and set to false `flag[i]` after completing the CS.
- Ti does not enter the CS while `flag[j]` is true.
- If only one task wants to enter the CS several times, it has no problems, as the other flag is always unset

Task T0:

```
while(true)
```

```
{
```

```
    flag[0] = true;
```

```
    while(flag[1]) { /*bw*/ }
```

Critical Section

```
    flag[0] = false;
```

Remainder Section

After  
you sir!

Task T1:

```
while(true)
```

```
{
```

```
    flag[1] = true;
```

```
    while(flag[0]) { /*bw*/ }
```

Critical Section

```
    flag[1] = false;
```

Remainder Section

After  
you sir!

# Algorithm 2

Discussion:

- Does it ensure mutual exclusion?
  - Yes, a process in CS has its flag set, but it does not enter CS if the flag of the other process is set
- Does it satisfy the bounded waiting requirement?
  - If the task  $T_i$  is scheduled after the  $T_j$  has set  $\text{flag}[j]$  to false but before it set again  $\text{flag}[j]$  to true,  $T_i$  will enter its CS.
- Does it satisfy the progress requirement?
  - What happens after the following sequence of instructions:
    - $T_0: \text{flag}[0] = \text{true};$
    - $T_1: \text{flag}[1] = \text{true};$
  - No progress, deadlock occurs.

Task  $T_i$ :

```
while(true)  
{
```

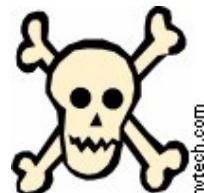
```
    flag[i] = true;  
    while(flag[j]) {/*bw*/ }
```

Critical Section

```
    flag[i] = false;
```

Remainder Section

```
}
```



ntech.com

# Algorithm 3

So, neither Algorithm 1, nor Algorithm 2 work. What do we do?

Use the ideas from both of them!

- Use the flag `i` to indicate willingness to enter CS
- But use `turn` to let the other task enter the CS

Task T0:

```
while(true)
{
```

```
    flag[0] = true;
    // T0 wants in
    turn = 1;
    // T0 gives a chance to T1
    while
        (flag[1]==true&&turn==1) {}
    Critical Section
    flag[0] = false;
    // T0 wants out
    Remainerd Section
}
```

Task T1:

```
while(true)
{
    flag[1] = true;
    // T1 wants in
    turn = 0;
    // T1 gives a chance to T0
    while
        (flag[0]==true&&turn==0) {}
    Critical section
    flag[1] = false;
    // T1 wants out
    Remainerd Section
}
```

# Algorithm 3

Lets see if it works ...

- Does it ensure mutual exclusion?
  - Yes, a task enters CS only if it is its turn when both want in.
- Does it satisfy bounded waiting?
  - Yes, when a task is in its RS, the other task can enter its CS (flag of other task is false).
- Ok, that was the easy part, lets go the crucial question:
- Does it satisfy the progress requirement?
  - A task will enter CS if and only if it is its turn OR the other task does not want to enter CS
  - OK, so if the other task does not want to enter, I can go in, great
  - But that was the easy part. What if both want to enter?

# Algorithm 3

Discussing progress requirement for Algorithm 3...

What happens if both tasks want to enter?

- Both will set their respective flags to true
- Both will set turn to the opposite value
  - Wait! Only one of them (the second one, for example T1) will really succeed in this, as turn can take only one value
- This means T0 will enter the CS!
- And after T0 exits it, it will set its flag to false and T1 will enter CS.
- Nonsense! There must be a way they block each other. Perhaps they start staggered...

# Both tasks want to enter!

Task T0:

```
CS  
flag[0] = false;  
// T0 wants out  
RS  
flag[0] = true;  
// T0 wants in  
turn = 1;  
// T0 gives a chance to T1  
// but T1 cancels this action  
  
while  
(flag[1]==true&&turn=1) {};  
// test false, enter  
Critical Section  
flag[0] = false
```

Task T1:

```
flag[1] = true;  
// T1 wants in  
turn = 0;  
// T1 gives a chance to T0  
  
while  
(flag[0]==true&&turn=0) {};  
// test true, must wait
```

# Algorithm 3

Task Ti:

```
while(true)
{
    flag[i]=true; // I want in
    turn=j;        // but I let the other in
    while (flag[j]==true && turn==j){} ;
    Critical section
    flag[i]=false; // I no longer want in
    Remainder section
}
```

- Let's construct a situation where both tasks are blocked...
- So, lets assume T0 is blocked.
- The only possible place is in the while loop, therefore flag[1] is true and turn = 1.
- However, if turn=1, then T1 cannot be blocked.
- It might actually work!
- It indeed works, the algorithm is called Peterson's algorithm

# A few additional ideas

- A solution to the CSP is robust relative to tasks failing in the RS.
  - But, if a task fails in the CS, the other task shall be blocked.
- The Peterson algorithm can be generalised to more than 2 tasks
  - But, there exists other more elegant algorithms – the baker's algorithm
- For tasks with shared variable to work properly, all threads involved must used the same synchronization algorithm
  - A common protocol

# Critique of software solutions

- Difficult to program! And to understand!
  - The subsequent solutions we shall study are based on the existence of specialized instructions, which makes the work easier.
- Threads/processes that desire entry into their CS are **busy waiting**; which consumes CPU time.
  - For long critical sections, it is preferable to **block** threads/processes that must wait....

# Hardware Solution: disable interrupts

Simple solution:

- A process would not be preempted in its CS

Process Pi:

```
while(true)
```

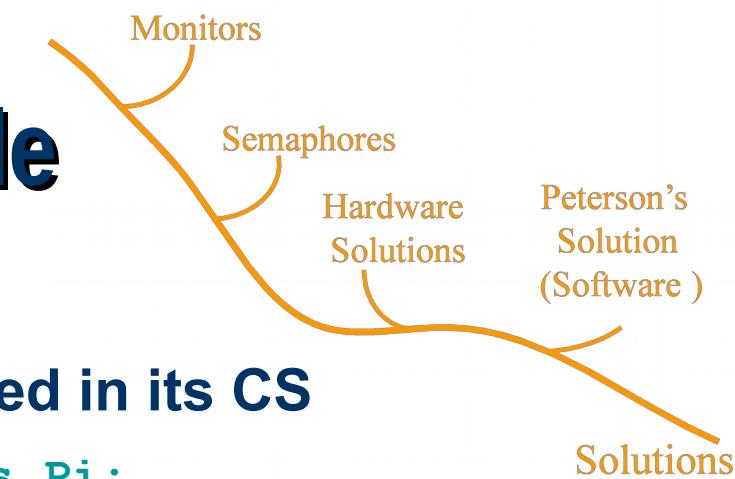
```
{
```

```
    disable interrupts  
    critical section  
    enable interrupts  
    remainder section
```

```
}
```

Discussion:

- Efficiency deteriorates: when a process is in its CS, it's impossible to interlace the execution of other processes in their RS.
- Loss of interruptions
- On a multiprocessor system: mutual exclusion is not assured.
- A solution that is not generally acceptable.



# Hardware solutions: specialized instructions

- Normal: when a thread or process access an address of memory, no other can access the same memory address at the same time.
- Extension: define machine instructions that perform **multiple** actions (ex: reading and writing) in the same memory location **atomically** (**indivisible**).
- An atomic instruction can only be **executed by a single thread/process** at a time (even in the presence of multiple processors).

# The test-and-set instruction

- Pseudocode for test-and-set:

```
bool testset(int &var)
{
    if (*var==0)
    {
        *var=1;
        return true;
    }
    else
    {
        return false;
    }
}
```

- An algorithm using test-and-set for mutual exclusion:
  - Variable b is initialized to 0
  - The first Ti to set b to 1 enters its CS.

Task Ti:

```
while testset(b)==false {};
CS //enter when true
b=0;
RS
```

Atomic instruction!

# The test-and-set instruction (cont.)

- Mutual exclusion is assured: if  $T_i$  enters the CS, the other  $T_j$  are busy waiting.
  - Problem: still using busy waiting.
- Can easily attain mutual exclusion, but needs more complex algorithms to satisfy the other requirements to the CSP.
  - When  $T_i$  leaves its CS, the selection of the next  $T_j$  is arbitrary: no bounded waiting: starvation is possible.

# Using xchg for mutual exclusion (Stallings)

- The variable **b** is initialized to 0
- Each Ti owns a *local variable k*
- The Ti that can enter its CS is the one that finds b=0
- The Ti excludes all others by assigning 1 to b
  - When the CS is occupied, both k and b will have the value 1 in tasks trying to enter its CS.
  - But k is 0 in the task that has entered its CS.
- xchg(a,b) suffers from the same problems as the test-and-set

- Certain CPUs (ex: Pentium) offer the instruction xchg(a,b) that exchanges atomically two variables.

usage:

```
Task Ti:  
while  
{  
    k = 1  
    while(k != 0) xchg(k,b);  
    Critical Section  
    xchg(k,b);  
    Remainder Section  
}
```

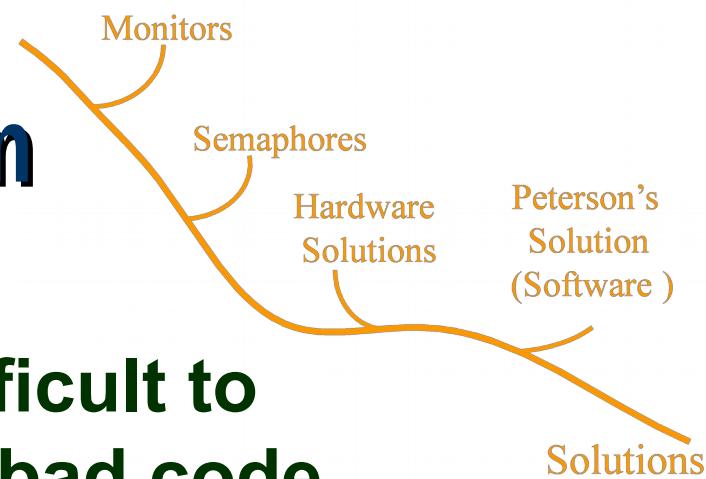
# **Drawback of software & hardware solutions**

- Processes that are requesting to enter in their critical section are **busy waiting** (consuming processor time needlessly)

Idea:

- If Critical Sections are long, it would be more efficient to block processes that are waiting...
- Can the OS help us? It can block processes.

# Solutions based on system calls – the semaphore



- **Solutions seen so far are difficult to program and easily leads to bad code.**
- **Need to find methods to avoid common errors, such as deadlocks, starvation, etc.**
  - **Need some higher level functions.**
- **The subsequent methods we shall study use more powerful functions that the OS can provide with system calls.**
- **Lets first study sempahores.**

# Semaphores – Version 1 – spinlocks

- The semaphore is a simple integer variable S on which three operations are allowed:
  - Initialization to a positive value (including zero).
  - **wait(S)**
    - Originally called P() (proberen → to test)
    - Also called acquire()
  - **signal(S)**
    - Originally called V() (verhogen → to increment)
    - Also called release()
- The first version of semaphores studied use busy wait.

# Spinlocks: Busy Wait Semaphores

- Easiest way to implement semaphores.
- Used in situations where waiting is brief or with multi-processors.
- When  $S=n$  ( $n>0$ ), up to  $n$  processes will not block when calling wait().
- When  $S$  becomes 0, processes block in the call wait() up until signal() is called.
- A call to signal() unblocks a blocked process or increments the semaphore value.

```
wait(S)
{
    while(S<=0) ;//no-op
    S--;
}
```

*The sequence  $S<=0$ ,  $S-$  must be atomic.*

```
signal(S)
{
    S++;
}
```

*The signal call must be atomic.*

# Using semaphores as a Solution to the CSP.

- Can be applied with n processes.
- Initialize S to 1.
  - Means only 1 process can enter its CS.
- To allow k processes to enter their CS, initialize S to k.

```
do
{
    wait(S);
        /*Critical Section*/
    signal(S);
        /*Remainder Section*/
} while(TRUE);
```

# Using semaphores for process synchronization

- We have 2 processes: P1 and P2
- Statement S1 in P1 needs to be performed before statement S2 in P2
- Can we do that using semaphores?

Thinking....

- P2 has to wait until P1 executes S1 ...
- ... then P1 needs to signal P2 that it can execute S2
- So, let P2 issue `wait(S)` before S2 and P1 do `signal(S)` after S1
- How should we initialize S?
  - `S.value = 1` does not work, as P2 can be scheduled first and execute statement S2
  - `S.value = 0` works

Process P1:

S1

`signal(S);`

Process P2:

`wait(S);`

S2

# Semaphores: observations

```
wait(S)
{
    while S<=0 {} ;
    S-- ;
}
```

- When  $S$  becomes  $>1$ , the process that unblocks first is the first to test  $S$  (random)
  - Possibility of famine exists.
  - This is not true for Version 2
- Still using busy waiting.
- With version 1 (spinlocks):  $S$  is always  $\geq 0$ .
- It is possible to limit the values of the semaphore to 0 and 1
  - This is the binary semaphore and also called the mutex (for mutual exclusion).
  - Easier to implement
- The more general semaphore is called the counting semaphore

# Semaphores – Version 2 – no busy wait

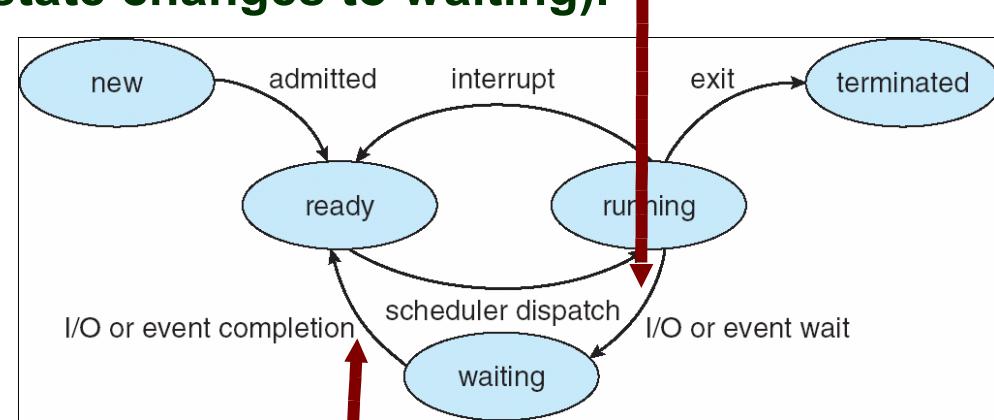
- When a process must wait for a semaphore to become greater than 0, place it in a queue of processes waiting on the same semaphore
- The queues can be FIFO, with priorities, etc. The OS controls the order in which processes are selected from the queue.
- Thus *wait* and *signal* become system calls similar to requests for I/O.
- There exists a queue for each semaphore similar to the queues defined for each I/O device.

# Semaphore – Version 2

- The semaphore S now becomes a data structure:

```
typedef struct  
{  
    int value;  
    struct process *list;  
} semaphore;
```

- A process waiting on a semaphore S, is blocked and placed in the queue S.list (its state changes to waiting).



- Signal(S) removes (according to the queuing discipline, such as FIFO) a process from S.list and places it on the ready queue (process enters the ready state).

# Implementation of wait()

```
wait(semaphore *S)
{
    S->value--;
    if(S->value < 0)
    {
        add the process to S->list;
        block(); /*process goes to state waiting*/
    }
}
```

- This system call must be atomic.
- When value becomes negative, its absolute value represents the number of blocked processes in list.
  - Recall that spinlock values never become negative (note that decrement and test operations are reversed)

# Implementation of signal

```
signal(semaphore *S)
{
    S->value++;
    if(S->value <= 0)
    { /* processes are waiting */
        Remove process P from S->list;
        wakeup(P); /*process becomes ready*/
    }
}
```

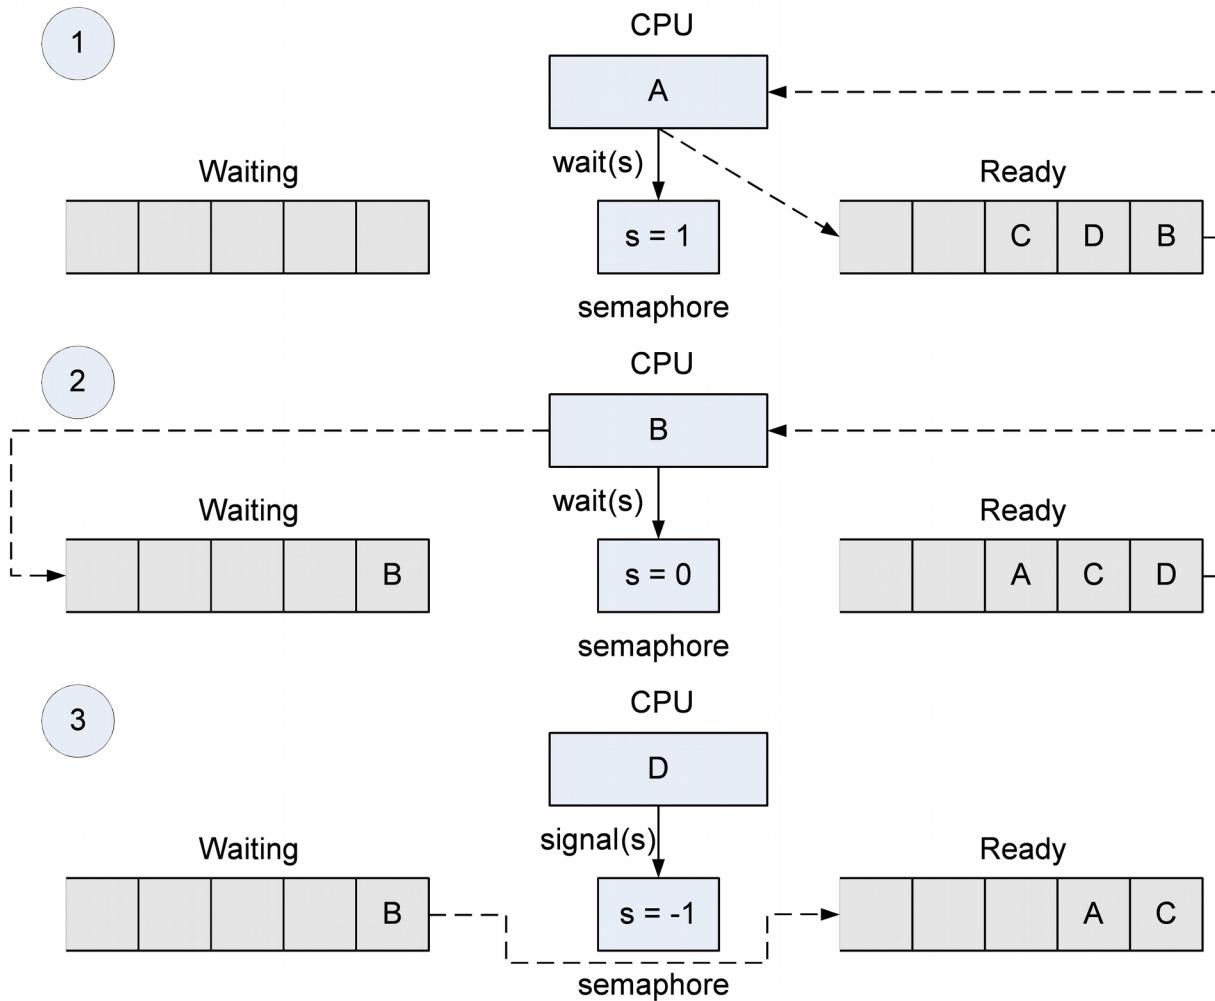
- This system call must be atomic
- The value of the semaphore can be initialised to a positive value (including 0).

# **Atomic execution of the wait() and signal() calls**

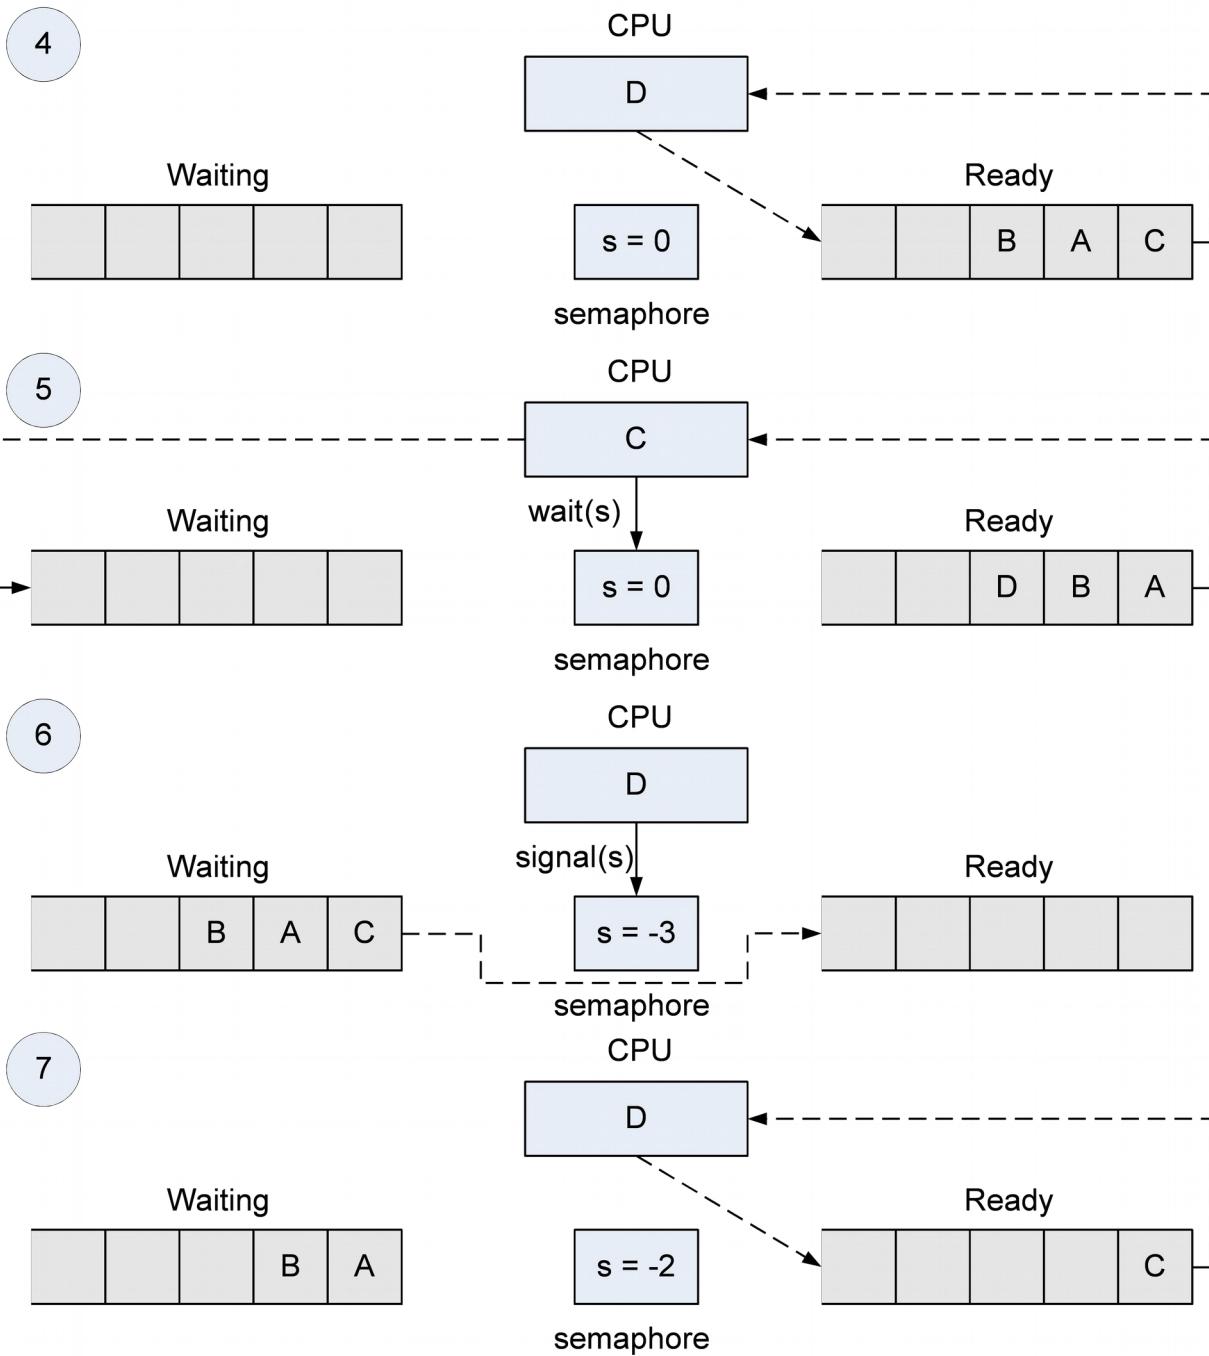
- wait() and signal() are in fact critical sections.
- Can we use semaphores for these critical sections?
- With single processor systems
  - Can disable interrupts
  - Operations are short (about 10 instructions)
- With SMP systems
  - Spinlocks
  - Other software and hardware CSP solutions.
- Result: we have not eliminated busy waiting
  - But the busy waiting has been reduced considerably (to the wait() and signal() calls)
  - The semaphores (version 2), without busy wait, are used within applications that can spend long periods in their critical section (or blocked on a semaphore waiting for a signal) – many minutes or even hours.
- Our synchronization solution is thus efficient.

# Example of semaphore operation (Stallings)

- **Process D**
  - Produces some result and uses signal(s) to indicate the availability of results to other processes
- **Processes A, B, C:**
  - Consumes the results from D, and uses wait(s) to access them.
- **Scenario 1:**
  - D has already indicated that a result is ready and thus  $s=1$

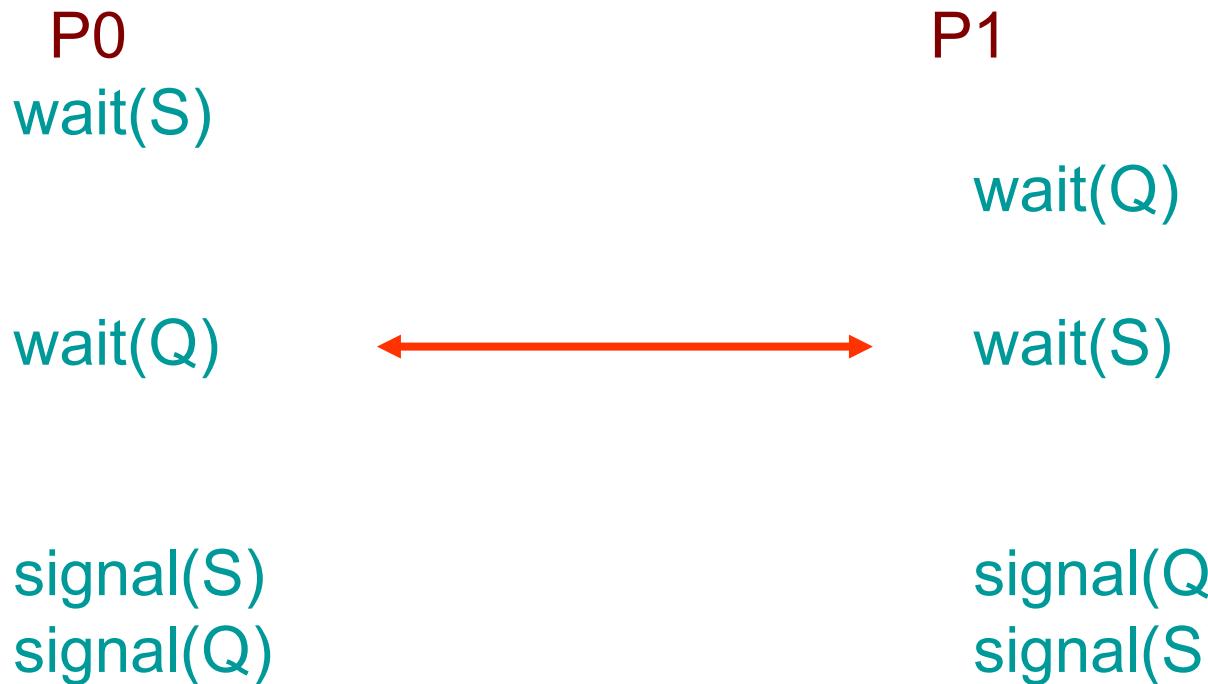


# Example of semaphore operation (Stallings)



# Deadlock and Starvation with Semaphores

- **Starvation:** consider what can happen if a LIFO queuing discipline were used in a semaphore.
- **Deadlock:** Given semaphores S and Q both initialized to 1



Reader  
Writer

Bounded  
Buffer

Dining  
Philosophers

We will come back  
to monitors later

Monitors

Semaphores

Hardware  
Solutions

Peterson's  
Solution  
(Software )

Solutions

## Process Synchronization

Classic  
problems

What is the  
problem  
?

In  
chi.  
Syn  
Real Systems

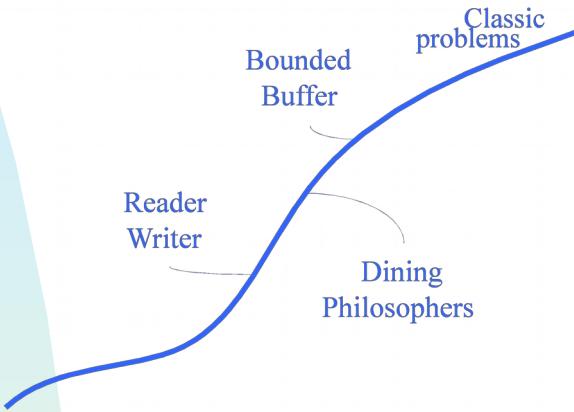
Background  
Producer /Consumer

Critical  
Section  
Problem

Mutual  
Exclusion

Bounded  
Waiting

# Classical Problems of Synchronization



- **Why to study them?**
- **These are abstractions of the typical synchronization problems that occur again and again**
  - Demonstrate the typical issues in process synchronization
- **Will allow us to practice using the semaphores**

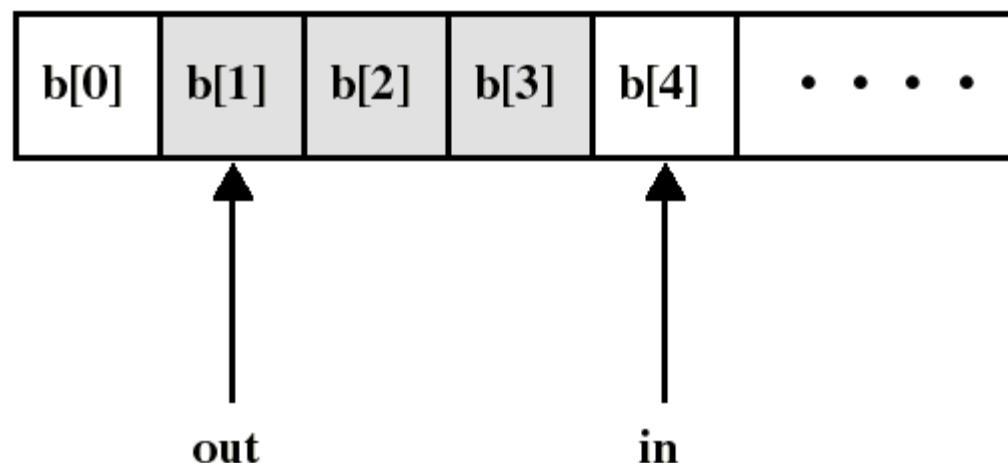
# The Bounded Buffer problem

- A **producer process** produces information that is consumed by a **consumer process**
  - Ex1: a print program produces characters that are consumed by a printer
  - Ex2: an assembler produces object modules that are consumed by a loader
- We need a **buffer** to hold items that are produced and eventually consumed
- A common paradigm for cooperating processes

# Start with an Unbounded buffer

- We assume first an **unbounded** buffer consisting of a linear array of elements
- **in** points to the next location in the buffer to receive the item to be produced
- **out** points to the next item to be consumed

shaded area indicates portion of buffer that is occupied



# Unbounded buffer

## How to start?

- We need to ensure mutual exclusion when accessing the buffer
  - Have a semaphore `bmutex` initialized to 1
  - Call `wait(bmutex)` before each access of the buffer
  - Call `signal(bmutex)` when finished accessing the buffer
- How to make sure the consumer does not consume something that was not yet produced?
  - Have another semaphore `num`, where `num.value` is the number of produced, but not yet consumed items
  - When producer produces an item, calls `signal(num)`
  - When consumer wants to consume item, calls `wait(num)`
  - How to initialize `num.value`?

# Solution for the unbounded buffer

Initialization:

```
bmutex.value = 1;  
num.value = 0;  
in = out = 0;
```

Producer:

```
while(true)  
{  
    v = produce();  
    wait(bmutex);  
    append(v);  
    signal(bmutex);  
    signal(num);  
}
```

Critical Sections

append(v)	take()
{	{
b[in] = v;	w = b[out];
int++;	out++;
}	return w;
	}

Consumer:

```
while(true)  
{  
    wait(bmutex);  
    w = take();  
    signal(bmutex);  
    consume(w);  
    wait(num);  
}
```

Does it actually work?

# Solution for the unbounded buffer

Initialization:

```
tmutex.value = 1;  
num.value = 0;  
in = out = 0;
```

Producer:

```
while(true)  
{  
    v = produce();  
    wait(bmutex);  
    append(v);  
    signal(bmutex);  
    signal(num);  
}
```

Critical Sections

```
append(v)      take()  
{  
    b[in] = v;      w = b[out];  
    in++;          out++;  
}  
return w;
```

Consumer:

```
while(true)  
{  
    wait(bmutex);  
    wait(num);  
    w = take();  
    signal(bmutex);  
    consume(w);  
}
```

Will it work now?

# Solution for the unbounded buffer

Initialization:

```
tmutex.value = 1;  
num.value = 0;  
in = out = 0;
```

Producer:

```
while(true)  
{  
    v = produce();  
    wait(bmutex);  
    append(v);  
    signal(bmutex);  
    signal(num);  
}
```

Critical Sections

```
append(v)      take()  
{  
    b[in] = v;      w = b[out];  
    in++;          out++;  
}  
return w;
```

Consumer:

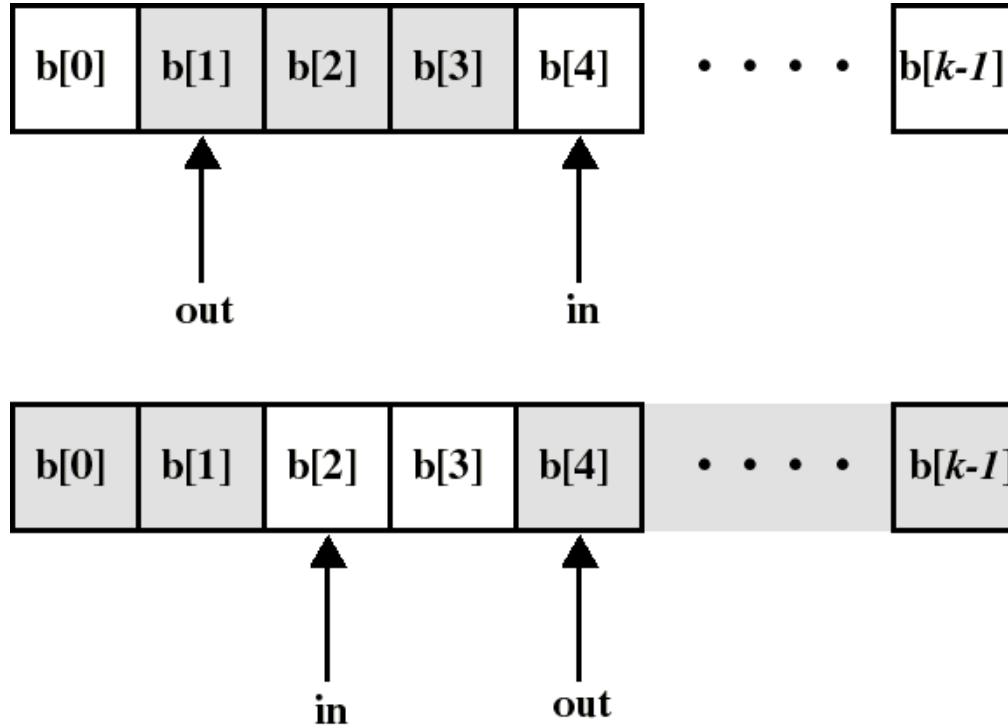
```
while(true)  
{  
    wait(num);  
    wait(bmutex);  
    w = take();  
    signal(bmutex);  
    consume(w);  
}
```

Will it work now? Yes!

# The Unbounded Buffer

- **Remarks:**
  - Putting `signal (num)` inside the CS of the producer (instead of outside) has no effect since the consumer must always wait for both semaphores before proceeding
  - The consumer must perform `wait (num)` before `wait (bmutex)` , otherwise **deadlock** occurs if consumer enter CS while the buffer is empty
- **Using semaphores is a difficult art...**

# Bounded buffer of size k



- can consume only when number `num.value` of (consumable) items is at least 1
- can produce only when number `empty.value` of empty spaces is at least 1

# Bounded Buffer

- As before:
  - we need a semaphore `bmutex` to have mutual exclusion on buffer access
  - we need a semaphore `num` to synchronize producer and consumer on the number of consumable items
    - don't consume if `num.value <=0`
- In addition:
  - we need a semaphore `emty` to synchronize producer and consumer on the number of empty spaces
    - don't produce if `emty.value <=0`
    - How to initialize `emty.value`?

# Solution for the bounded buffer

Initialization:

```
bmutex.value = 1;  
num.value = 0;  
emty.value = k;  
in = out = 0;
```

Producer:

```
while(true)  
{  
    v = produce();  
    wait(emty);  
    wait(bmutex);  
    append(v);  
    signal(bmutex);  
    signal(num);  
}
```

Critical Sections

```
append(v)           take()  
{  
    b[in] = v;          w = b[out];  
    in=(in+1)%k;      out=(out+1)%k;  
}  
return w;
```

Consumer:

```
while(true)  
{  
    wait(num);  
    wait(bmutex);  
    w = take();  
    signal(bmutex);  
    signal(emty);  
    consume(w);  
}
```

# Readers-Writers Problem

There are two types of processes accessing a shared database

- readers only reads the data, but does not modify them
- writers that want to modify the data

In order to maintain consistency, as well as efficiency, the following rules are used

- Several readers can access the database simultaneously
- A writer needs to have an exclusive access to the database (has a critical section)
  - No readers or other writers are allowed while a writer is writing
- What to do if there are several readers in the system and a writer arrives?
  - While there is a reader active, do not have new readers wait (first readers-writers problem).
  - No new reader is admitted if there is a writer waiting (second readers-writers problem).
  - Both might lead to starvation

# Readers-Writers Problem

Hmmm, how to start?

- Have a semaphore wrt for allowing only one write process to modify the database
  - The writer process can simply use it, i.e.
    - wait(wrt), followed by CS, and then signal(wrt)
- How to prevent a writer entering when there is a reader?
  - The first reader grabs the semaphore wrt, i.e. executes wait(wrt) to prevent the writer from entering its CS
  - The last reader releases semaphore wrt (i.e. executed signal(wrt))
  - Now the writer can enter
- We need a counter readers to know how many readers are in the system, to know when to grab and release the wrt semaphore
  - Only the first one and the last one should do that
- But then we need a second semaphore rmutex to guard access to readers

# Readers-Writers Problem

Initialization:

```
wrt.value = 1;  
rmutex.value = 1;  
readers = 0;
```

Writer:

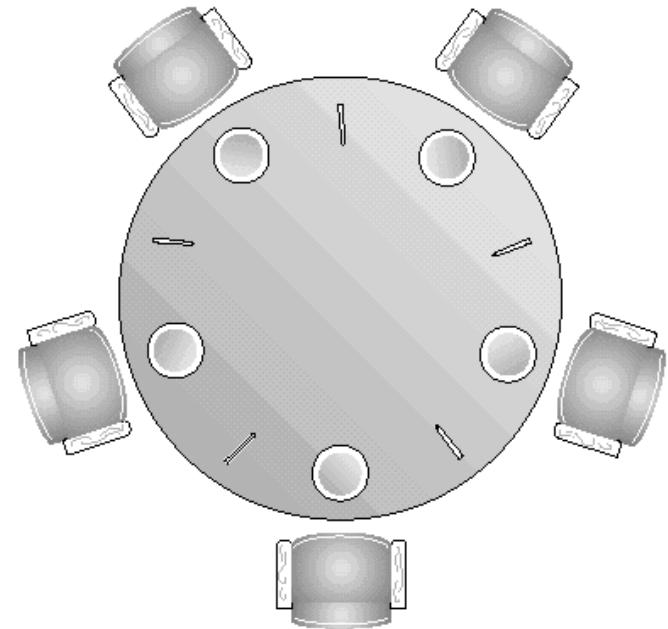
```
while(true)  
{  
    wait(wrt);  
    writer's CS;  
    signal(wrt);  
}
```

Reader:

```
while  
{  
    wait(rmutex);  
    readers++;  
    if (readers == 1)  
        wait(wrt);  
    signal(rmutex);  
    reading database;  
    wait(rmutex);  
    readers--;  
    if (readers == 0)  
        signal(wrt);  
    signal(rmutex);  
}
```

# The Dining Philosophers Problem

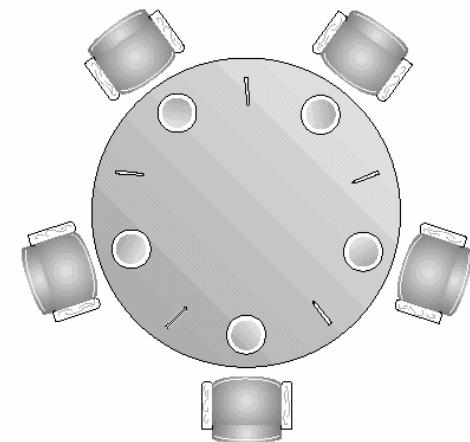
- 5 philosophers think, eat, think, eat, think ....
- In the center a bowl of rice.
- Only 5 chopsticks available.
- Require 2 chopsticks for eating.
- A classical synchronization problem
- Illustrates the difficulty of allocating resources among process without deadlock and starvation



# The Dining Philosophers Problem

- Each philosopher is a process
- One semaphore per chopstick:
  - So that two philosophers cannot grab the same chopstick simultaneously
  - semaphore chopst[5];
  - Initialization:  
chopst[i].value=1 for i = 0 to 4

```
Process Pi:  
while(true)  
{  
    think();  
    wait(chopst[i]);  
    wait(chopst[(i+1)%5]);  
    eat();  
    signal(chopst[i+1%5]);  
    signal(chopst[i]);  
}
```



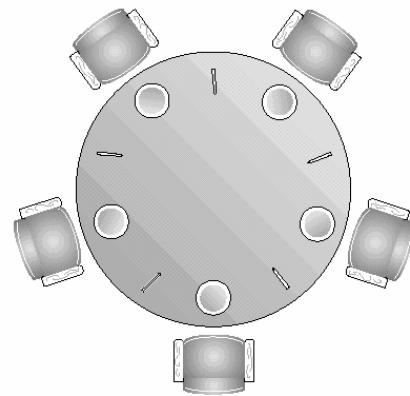
## Does it work?

- What happens if each philosopher starts by picking his left chopstick?
- Deadlock! Poor philosophers, they should know better...

# The Dining Philosophers Problem

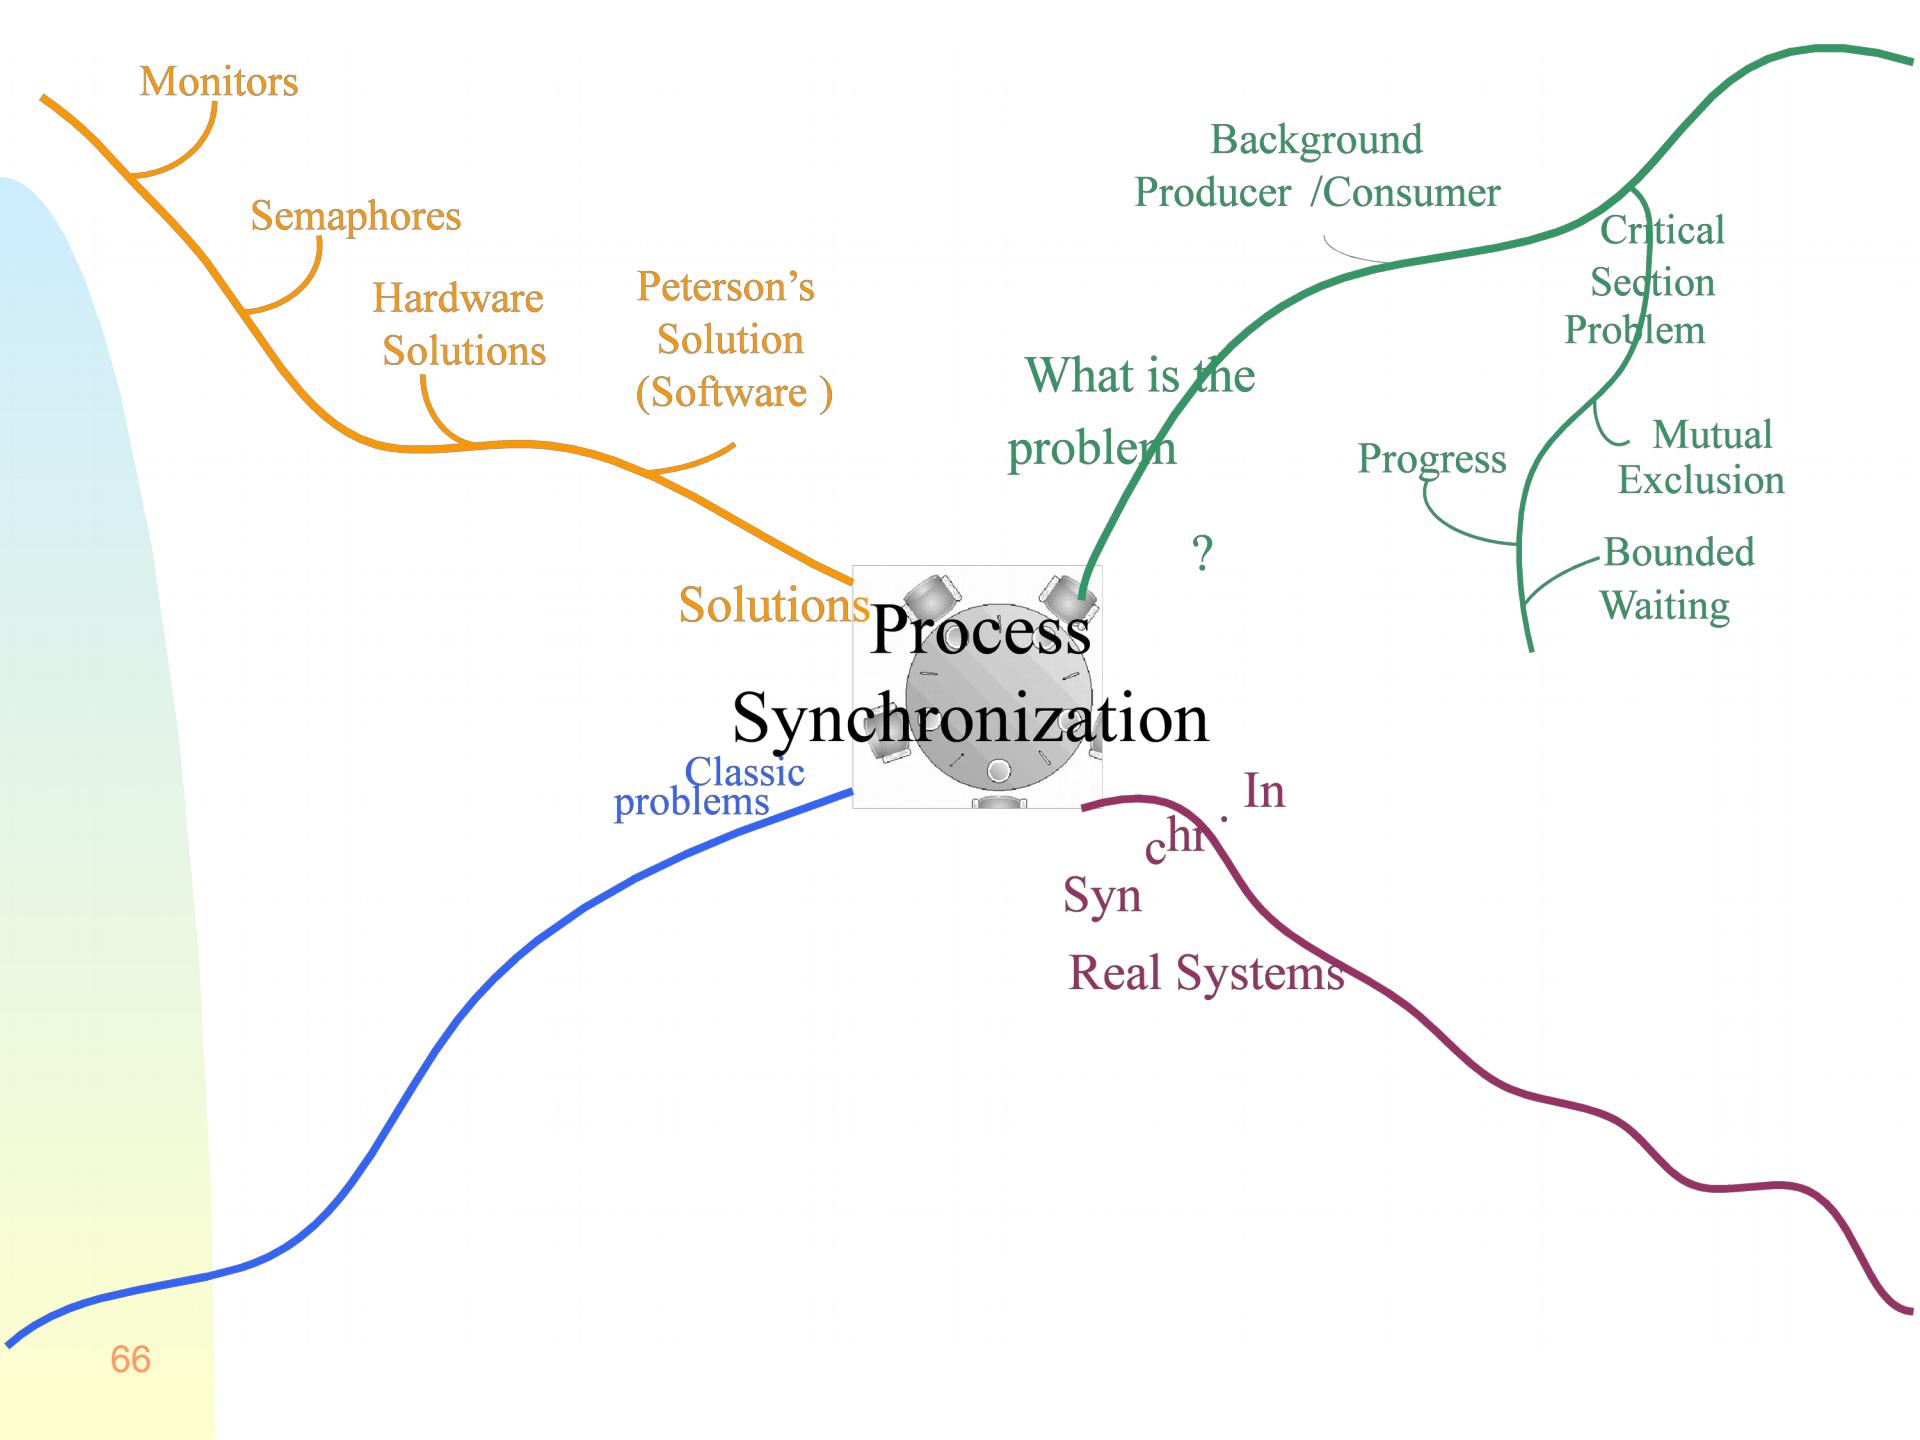
- A solution: allow only 4 philosophers at a time to try to eat
- Then 1 philosopher can always eat when the other 3 are holding 1 chopstick
- Hence, we can use another semaphore T that would limit at 4 the number of philosophers trying to take a chopstick
- How to initialize T.count?
  - T.value = 4
- Solution is free of deadlock and starvation

```
Process Pi:  
while(true)  
{  
    think();  
    wait(T);  
    wait(chopst[i]);  
    wait(chopst[(i+1)%5]);  
    eat();  
    signal(chopst[(i+1)%5]);  
    signal(chopst[i]);  
    signal(T);  
}
```



# **Advantages of semaphores (relative to other synchr. solutions)**

- **Single variable (data structure) per critical section.**
- **Two operations: wait, signal**
- **Can be applied to more than 2 processes.**
- **Can have more than a single process enter the CS.**
- **Can be used for general synchronization.**
- **Service offered by OS, including blocking and queue management.**



# Problems with semaphores: programming difficulty

- Wait and signal are scattered across different programs and running threads/processes
  - Not always easy to understand the logic
- Usage must be correct in all threads/processes
- One « bad » thread/process can make a collection of threads/processes fail (e.g. forget a signal)

# Monitors

- Are high-level language constructs that provide equivalent functionality to that of semaphores but are easier to control
- Found in many concurrent programming languages
  - Concurrent Pascal, Modula-3, C#, Java...
- Can be implemented with semaphores...
  - See section 6.7.3 of the textbook

# Monitors

- Is a software module (ADT – abstract data type) containing:
  - one or more procedures
  - an initialization sequence
  - local data variables
- Characteristics:
  - local variables accessible only by monitor's procedures
  - a process enters the monitor by invoking one of it's procedures
  - **only one process can be executing in the monitor at any one time** (but a number of threads/processes can be waiting in the monitor).

# Monitors

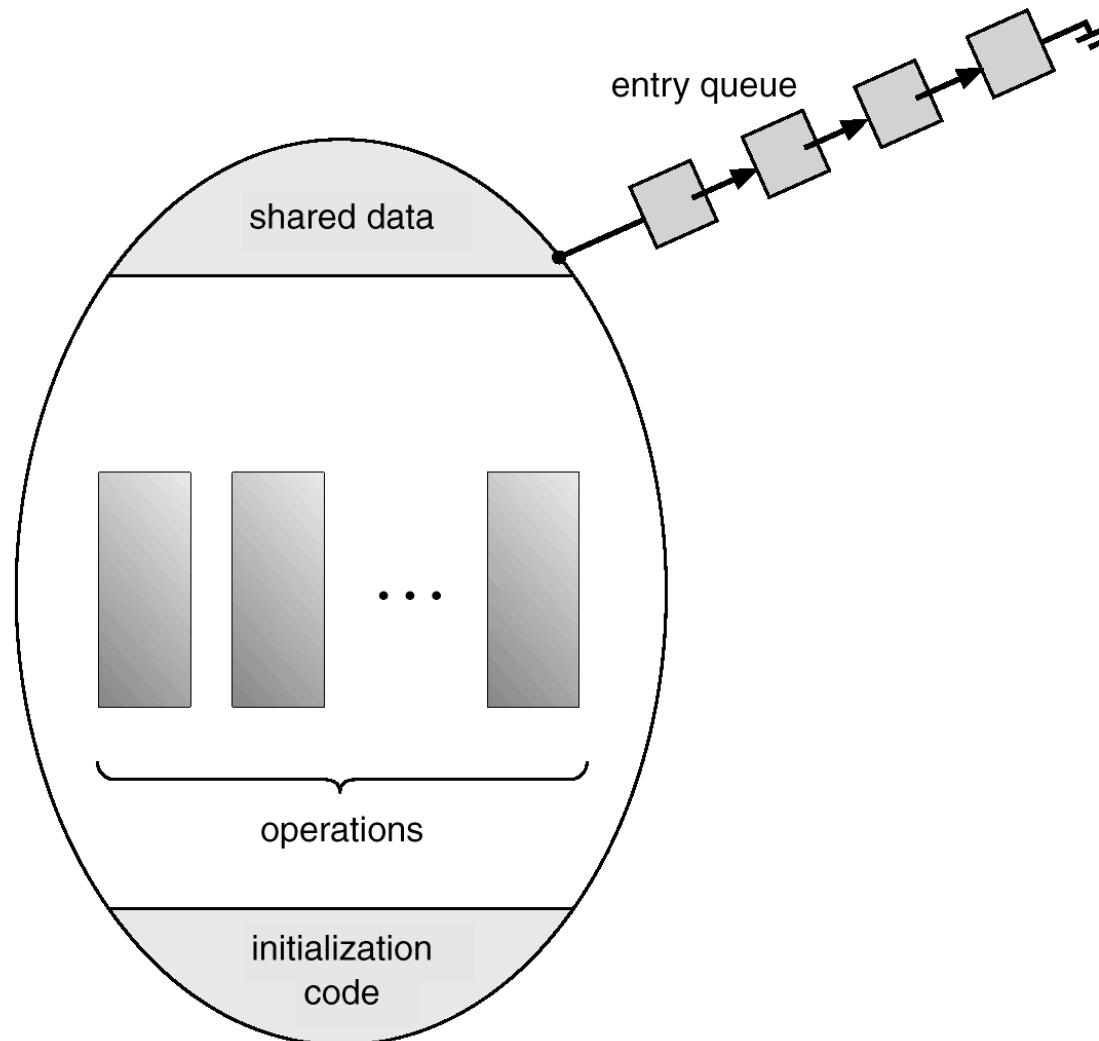
- The monitor ensures mutual exclusion: no need to program this constraint explicitly
- Hence, shared data are protected by placing them in the monitor
  - The monitor **locks** the shared data on process entry
- Process synchronization is done by the programmer by using **condition variables** that represent conditions a process may need to wait for before executing in the monitor

# General Structure of the Monitor

```
monitor monitor-name  
{  
    /* shared variable declarations */  
  
    p1(. . .) {p1 code}  
  
    p2(. . .) {p2 code}  
  
    . . .  
  
    initialization(. . .)  
        {initialization code}  
}
```

The only way to manipulate the shared internal variables in the monitor is by calling one of the procedures/functions/methodes p1, p2, ...

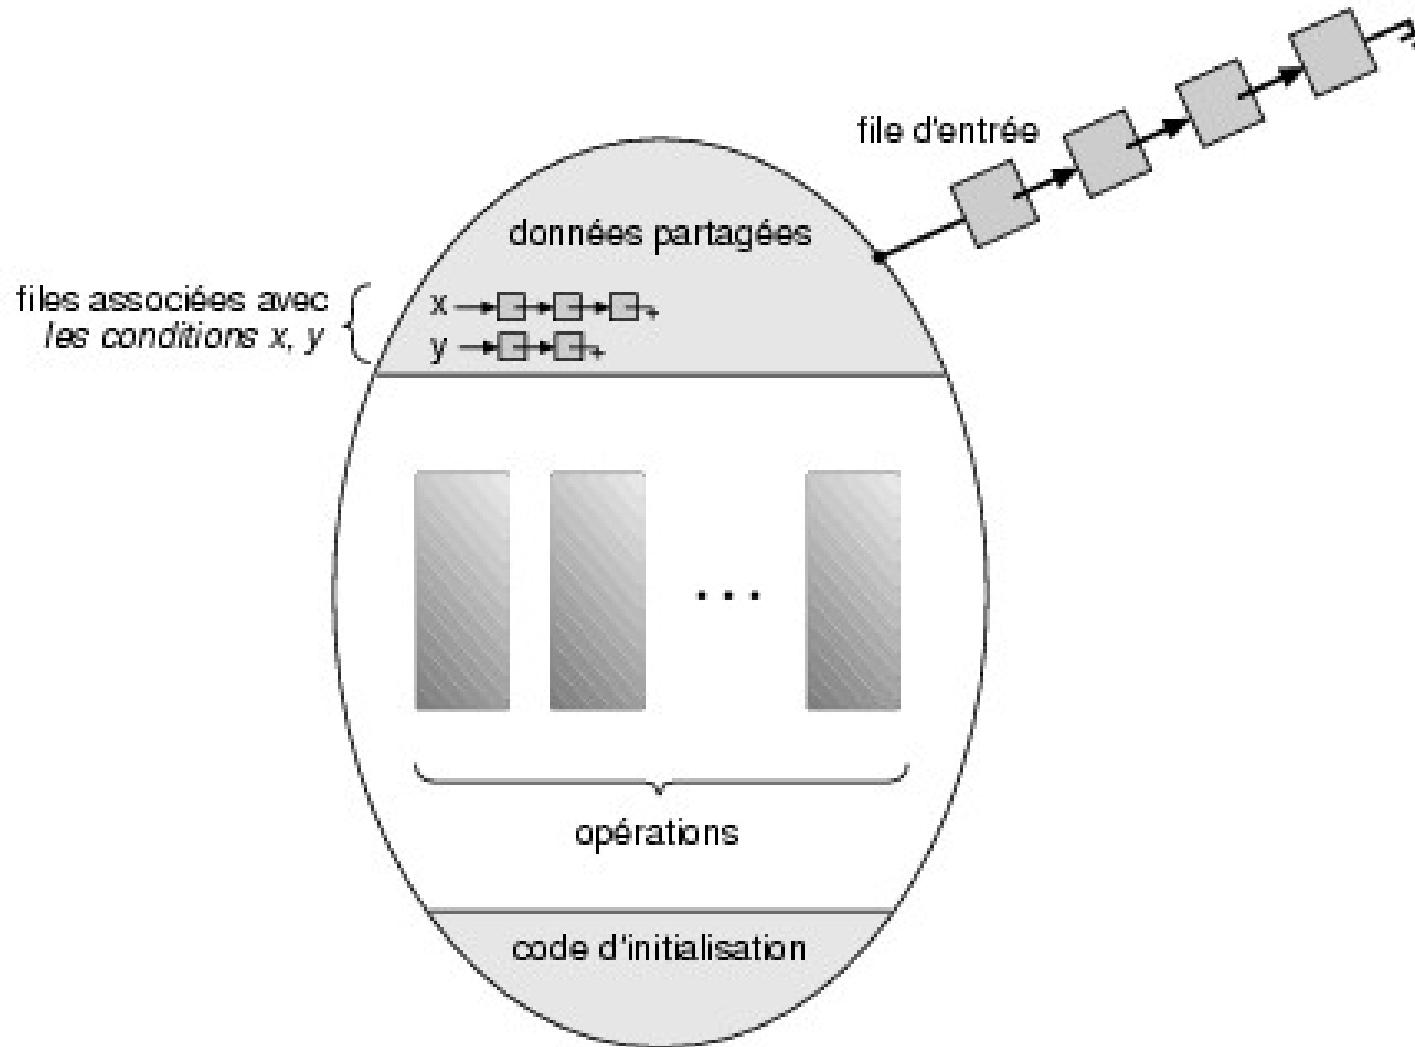
# Monitor: schematic view



# Condition variables

- are local to the monitor (accessible only within the monitor)
- can be accessed and changed only by two functions:
  - **x.wait()** blocks execution of the calling process/thread on condition (variable) *x*.
    - the process/thread can resume execution only if another process/thread executes *x.signal()*
  - **x.signal()** resume execution of some process/thread blocked on condition (variable) *x*.
    - If several such processes exist: choose one based on queuing scheme (e.g. FCFS, priority - see section 6.7.4 of the textbook)
    - If no such process exists: do nothing

# Monitor with Conditional Variables

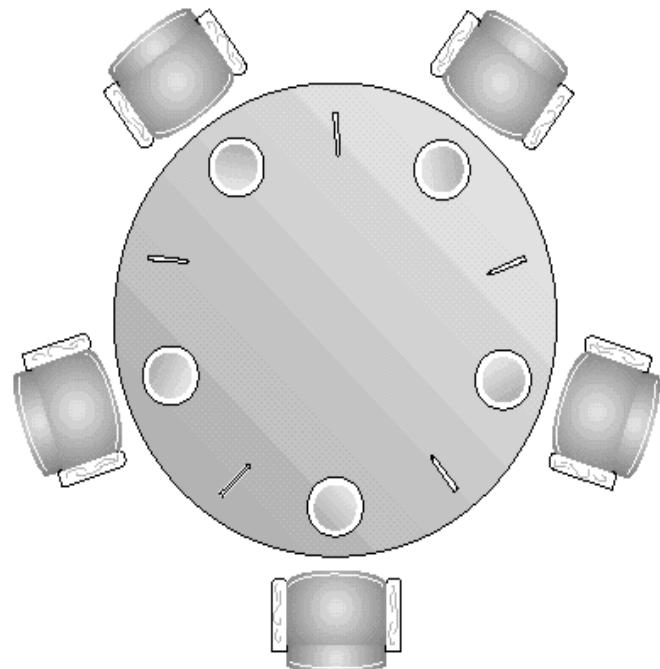


# A Problem with signal()

- Say a thread/process P executes x.signal() and frees a thread/process Q,
  - Now, 2 threads/processes that want to execute, P and Q – not allowed
- Two possible solutions:
  - Signal and Wait
    - P waits until Q leaves the monitor (e.g. in a special queue – *urgent* – see Stallings).
  - Signal and Continue
    - P continues its execution and Q waits until P leaves the monitor

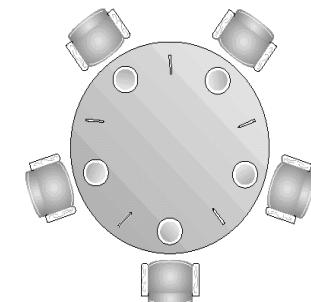
# **Back to the Dining Philosophers Problem**

- **5 philosophers think, eat, think, eat, think ....**
- **In the center a bowl of rice.**
- **Only 5 chopsticks available.**
- **Require 2 chopsticks for eating.**
- **A classical synchronization problem**



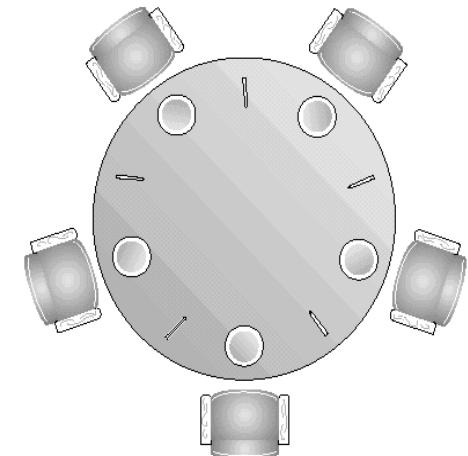
# Dinning Philosophers: A Monitor

- Define the monitor « dp » (dining philosophers)
- Shared variables: each philos. has its own state that can be: (thinking, hungry, eating)
  - philosopher i can only have state[i] = eating iff its neighbors are not eating
- Conditional variables: each philosopher has a condition self
  - the philosopher i can wait on self [ i ] if it wants to eat but cannot obtain 2 chopsticks
- Four functions in the monitor
  - pickup(i) – philo. i wants to pick up chopsticks
  - putdown(i) – philo. i puts down his/her chopsticks
  - test(i) – test the state of philo. i
  - Initialization\_code() - initialization



# Philosopher $i$ executes the loop:

```
while(true)
{
    think
    dp.pickup(i)
    eat
    dp.putdown(i)
}
```



# A philosopher eats

```
void test(int i)
```

```
{
```

```
    if ( (state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) )
```

```
{
```

```
    state[i] = EATING;
    self[i].signal();
```

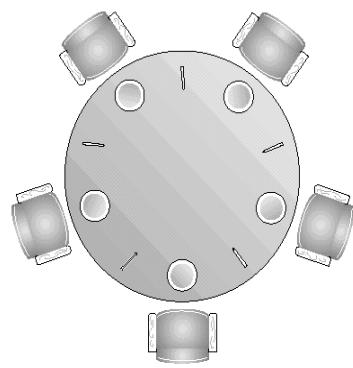
```
}
```

```
}
```

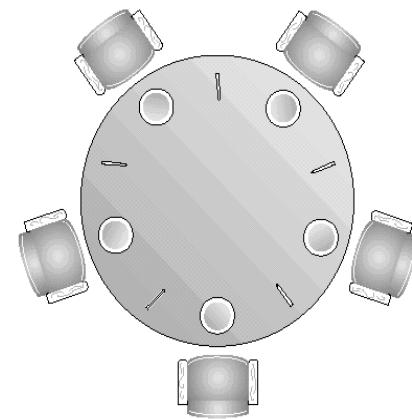
A philosopher eats if his/her neighbors do not eat and he/she is hungry.

When the philosopher starts eating, a signal allows the philosopher to leave the semaphore.

No change in state occurs if tests fail.



# Pickup Chopsticks

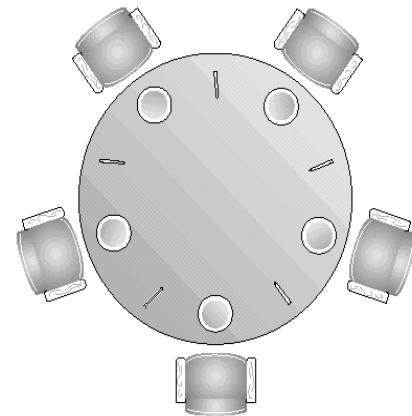


```
void pickUp(int i)
{
    state[i] = HUNGRY;
    test(i);
    if (state[i] != EATING)
        self[i].wait();
}
```

Phil. tries to start eating by doing a test. If comes out of test not eating, must wait – blocs until the execution of `self[i].signal()`

# Put down Chopsticks

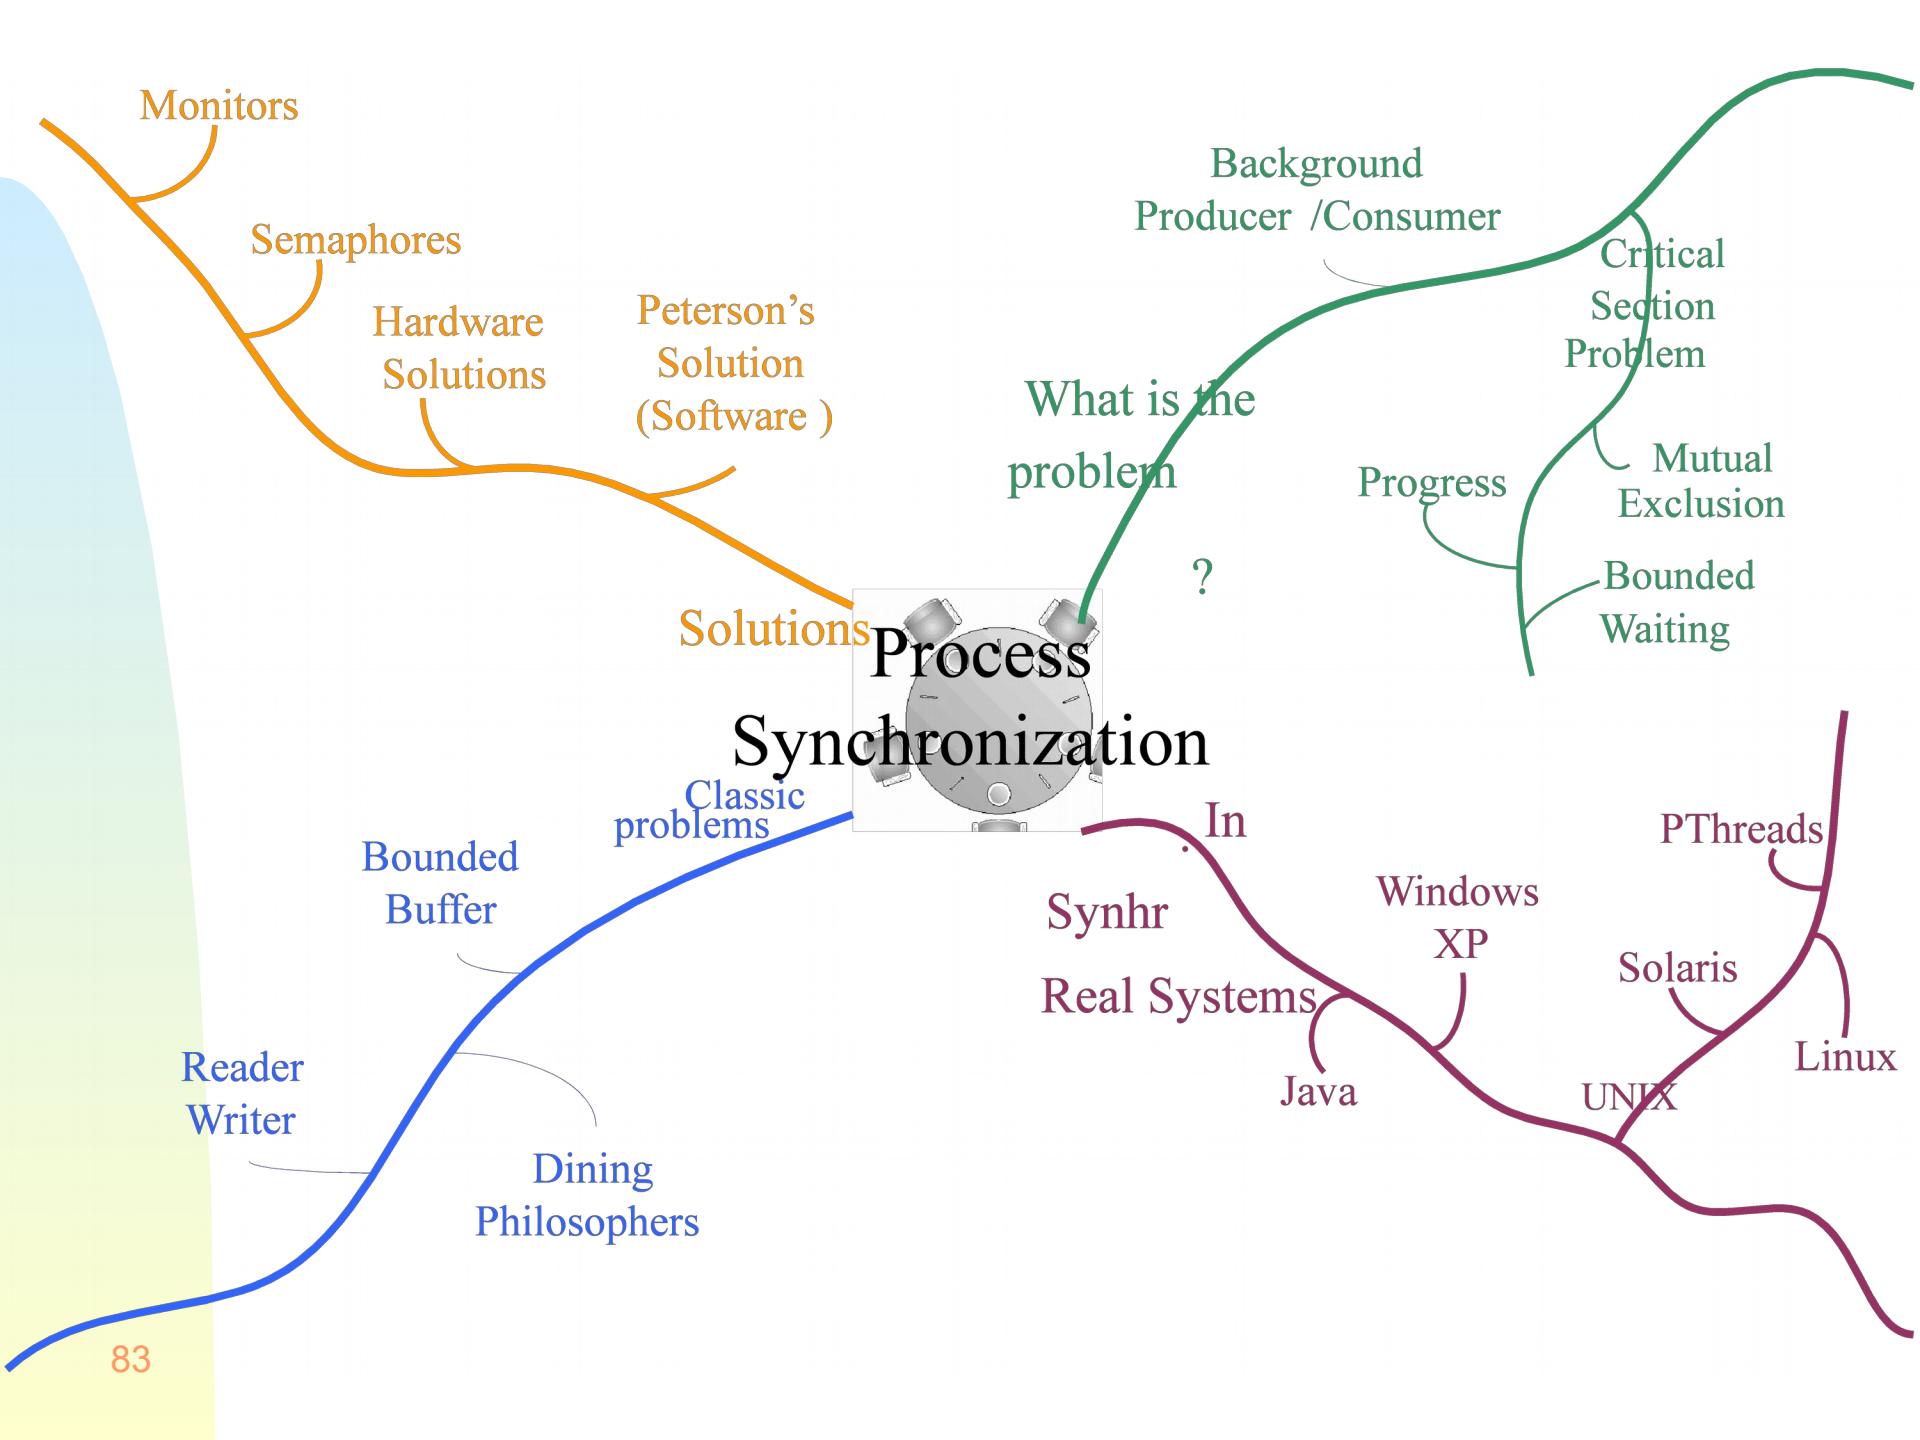
```
void putDown(int i)  
{  
    state[i] = THINKING;  
    // test the two neighbors  
    test((i + 4) % 5);  
    test((i + 1) % 5);  
}
```



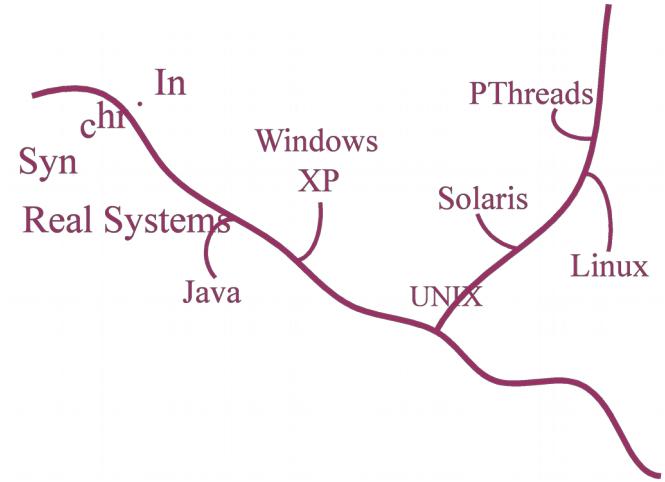
Once eating is done, a philosopher tries to get his/her neighbors eating using test.

**An ingenious solution – note that  
the chopsticks are implicit.**

**Worth studying.**



# Java Class Semaphore



- **Counting semaphore**
- **Characteristics :**
  - Can choose to use FIFO or not.
  - Semaphore value represents a number of available permits
  - `release()` increments this value – adds permits.
  - `acquire()` blocks if no permit is available.
  - In principle the number of permits is never negative.
    - One exception, it is possible to initialize the number of permits to a negative value.

# Semaphore Constructors

```
Semaphore(int permits)
```

```
Semaphore(int permits, boolean fair)
```

- **The parameter `permits`: defines the initial number of permits.**
- **The parameter `fair`: If `fair` is true the FIFO discipline is used with the semaphore queue.**
- **The first constructor gives a value of `false` to the parameter `fair`, that is, does not use FIFO with the queue.**

# The method `acquire()`

- This method plays the same role as the function `wait()` studied in class.
- Verifies the number of available permits in the semaphore.
  - If a permit is available (semaphore value > 0), the number of semaphores is decremented and the method returns; the thread can continue executing.
  - If no permit is available (semaphore value <= 0), the thread blocks and is placed on a queue.
    - When the permit becomes available, a thread in the queue is removed and allowed to consume the permit.

## The method `acquire()`

- This method will throw an exception when a thread is interrupted (with the method `interrupt()`).
  - It throws `InterruptedException`.
  - Use the method in a try/catch structure:

```
try{sem.acquire();}  
catch (InterruptedException e) { }
```
- Add `break;` as shown below to break out of a loop.

```
try{sem.acquire();}  
catch (InterruptedException e) { break; }
```

## The method `acquireUninterruptibly()`

- This method provides the same functionality as `acquire()`,
- But it does not throw an exception with the thread is interrupted.
- Not necessary to use the `try/catch` with this method.
- Use the method with threads that are not to be interrupted.

# The method `release()`

- This method plays the same role as the function `signal()`.
  - It increments the number of permits in the semaphore (i.e. its value).
  - If threads blocked on the semaphore, releases a thread that is blocked on the semaphore to consume the permit

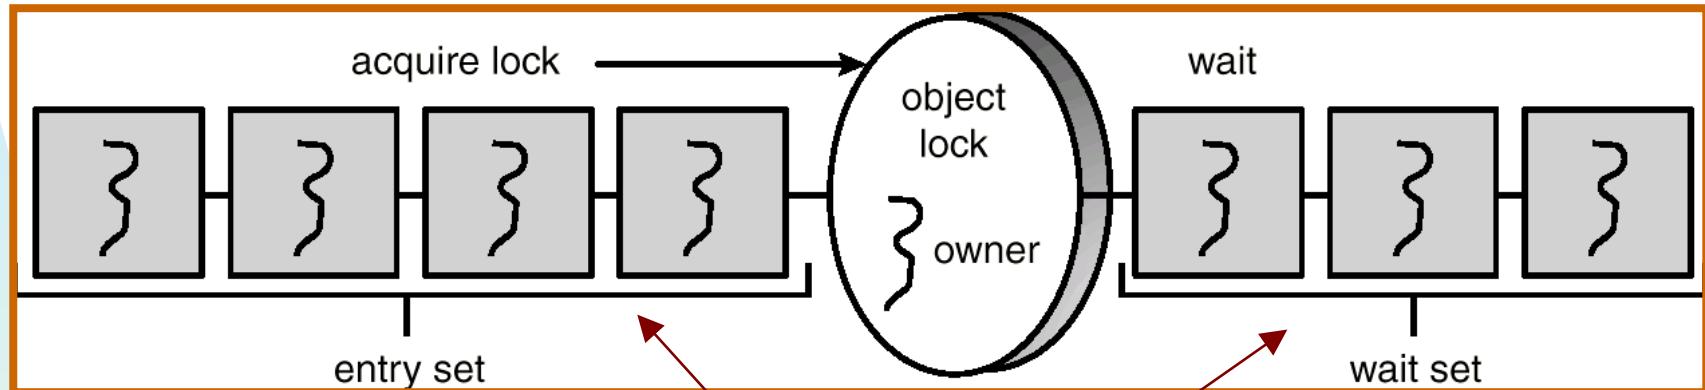
# Other methods in Semaphore

- See Java documentation for more details on other methods available in Semaphore.
  - Other versions of `acquire()` :  
`acquireUninterruptibly()` , `tryAcquire()` , and  
`release()`
  - `availablePermits()`
  - `drainPermits()` and `reducePermits()`
  - `isFair()`
  - `hasQueuedThreads()` , `getQueueLength()` ,  
`getQueuedThreads()`

# A simple monitor in Java

- Objects with synchronized methods in Java are essentially monitors (although simple)
  - A single thread at a time can execute such methods.
- Two queues are defined for each Java object:
  - Entry set
  - Wait set
- An object can only have a single wait set (queue)
  - Important limitation that can complicate things in Java...
- Wait in Java is more or less the same as wait() described for monitors
- Signal is called notify
  - Notify() frees 1 single thread
  - NotifyAll frees all threads
  - Freed threads are not executed: they are placed in the entry set.

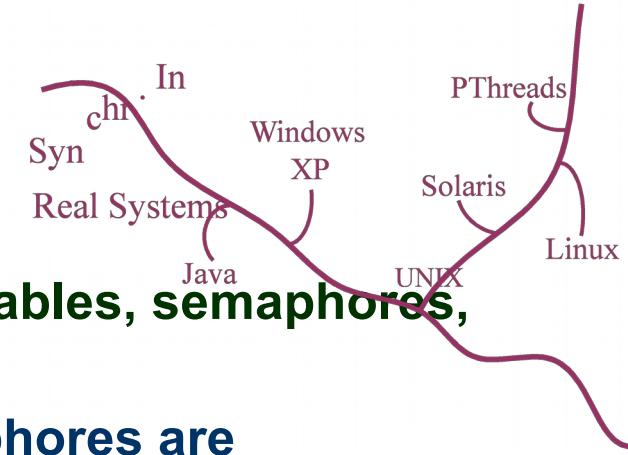
# Java Entry Sets and Wait Sets



```
public class SimpleClass
{
    ...
    public synchronized void SynchrMethod()
    {
        ...
        wait();
        ...
    }
}
```

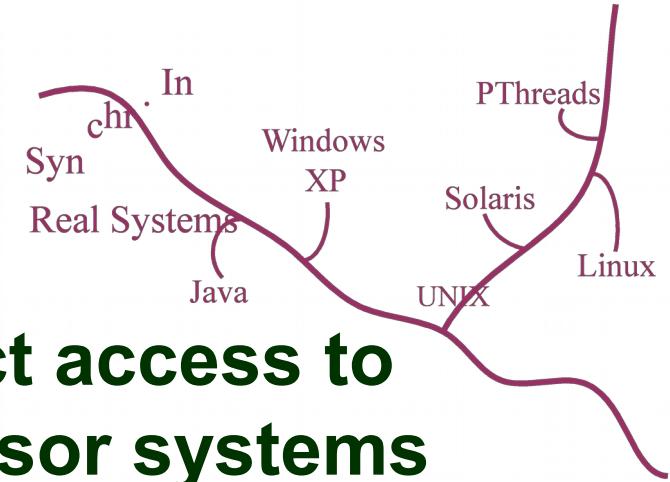
# Solaris Synchronization

- Implements adaptive mutexes, conditional variables, semaphores, reader-writer locks and turnstiles
  - Conditional variables (monitors) and semaphores are implemented as presented in this module
- Uses adaptive mutexes for efficiency when protecting data from short code segments
  - Starts as a spinlock if thread holding lock is likely to finish
  - Blocks if thread holding lock is not running
- Uses readers-writers locks for accessing data frequently for read only purposes
  - Expensive to implement – used on long sections of code.
- Uses turnstiles to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock
  - Queues associated to threads rather than synchr. objects
  - When a thread is the first object to block on an object, its turnstile becomes the blocking queue for the object

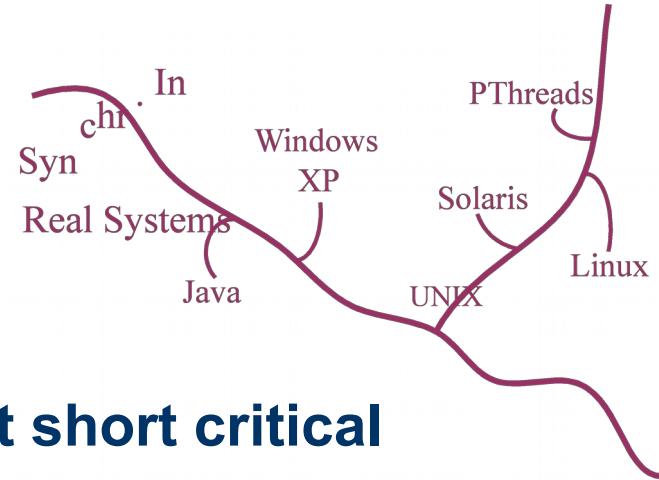


# Windows XP Synchronization

- Uses interrupt masks to protect access to global resources on uniprocessor systems
  - Uses spinlocks on multiprocessor systems
  - Threads holding a spinlock is never preempted.
- Also provides dispatcher objects for synchronization outside the kernel
  - Provides different mechanisms including mutexes, semaphores, events and timers.
  - Event: acts much like a condition variable

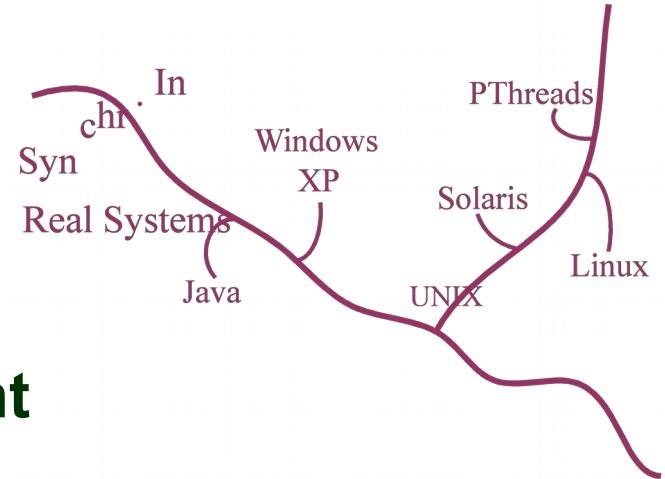


# Linux Synchronization



- On single processor systems:
  - disables interrupts to implement short critical sections
- On SMP systems
  - The spinlock is the fundamental locking mechanism.
  - Spinlocks/disabling interrupts are mechanisms used for holding lock during short periods.
  - Linux also provides semaphores:
    - Also reader-writer version is available.

# Pthreads Synchronization



- **Pthreads API is OS-independent**
- **It provides:**
  - **mutex locks**
  - **condition variables**
  - **read-write locks**
- **Non-portable extensions include:**
  - **semaphores**
  - **spin locks**

